

■ MIPS: Millions of Instructions Per Second

- Doesn't account for
 - Differences in ISAs between computers
 - Differences in complexity between instructions

$$\begin{aligned} \text{MIPS} &= \frac{\text{Instruction count}}{\text{Execution time} \times 10^6} \\ &= \frac{\text{Instruction count}}{\frac{\text{Instruction count} \times \text{CPI}}{\text{Clock rate}}} \times 10^6 = \frac{\text{Clock rate}}{\text{CPI} \times 10^6} \end{aligned}$$

- CPI is the average for a program
- and yields the “*native MIPS*”

- MIPS can be misleading
 - Varies between programs on the same machine
 - Does not reflect the instruction set complexity
 - May vary inversely with performance

- The following example illustrates this last point:

Assume a program is translated by two different compilers for the same machine:

$$\text{MIPS} = \frac{\text{Clock rate}}{\text{CPI} \times 10^6}$$

	Instruction counts (in millions) for each instruction class		
Instruction Class	A	B	C
Code from compiler 1	4	2	1
Code from compiler 2	9	1	1

$$\text{CPI}_1 = ((4*1 + 2*2 + 1*3)*10^6) / ((4+2+1)*10^6) = 1.57$$

$$\text{Thus MIPS}_1 = 100 \text{ MHz} / 1.57*10^6 = 63.64$$

$$\text{CPU time}_1 = ((4+2+1)*10^6* 1.57) / 100 * 10^6 = 0.11 \text{ sec}$$

Assume a program is translated by two different compilers for the same machine:

Instruction Class	Instruction counts (in millions) for each instruction class		
	A	B	C
Code from compiler 1	4	2	1
Code from compiler 2	9	1	1

$$CPI_2 = ((9*1 + 1*2 + 1*3)*10^6) / ((9+1+1)*10^6) = 1.27$$

$$\text{Thus } MIPS_2 = 100 \text{ MHz} / 1.27*10^6 = 78.57$$

$$CPU \text{ time}_2 = ((9+1+1)*10^6* 1.27) / 100 * 10^6 = 0.14 \text{ sec}$$

Yields a higher MIPS rating but takes longer to execute

- *Native MIPS* was defined above
- *Peak MIPS* is based on minimum CPI
 - Assumes all instructions in program have minimum CPI
 - Minimum possible CPI = 1
 - All instructions require at least 1 clock cycle
- Relative MIPS is based on a standard machine

$$MIPS_{relative} = \frac{Time_{reference}}{Time_{unrated}} \times MIPS_{reference}$$



- *Benchmarks are suites of programs*
 - Chosen to be typical of workloads
 - Systems are rated based on benchmark execution times
 - Average time is used as a measure of performance

- Arithmetic mean is not used as average
 - Outliers can have undue influence on result
 - Geometric mean is used instead (defined below)

Geometric Mean




$$\sqrt[n]{\prod_{i=1}^n \text{Execution time}}$$

- Standard Performance Evaluation Corp (SPEC)
 - Develops benchmarks for CPU, I/O, Web, ...
- SPEC CPU2006
 - Elapsed time to execute a selection of programs
 - Negligible I/O, so focuses on CPU performance
 - Normalize relative to reference machine
 - Summarize as geometric mean of performance ratios
 - CINT2006 (integer) and CFP2006 (floating-point)

- SPEC CPU2017
 - Comprised of a group of 43 programs
 - Organized into 4 suites

- Another possibility is the Harmonic Mean
 - More appropriate for averaging rates such as MIPS
 - Average MIPS for suite gives a single performance measure

Harmonic Mean 
$$\overline{S}_H = \frac{N}{\sum_1^N \frac{1}{S_i}}$$

- Improving an aspect of a computer and expecting a proportional improvement in overall performance

$$T_{\text{improved}} = \frac{T_{\text{affected}}}{\text{improvement factor}} + T_{\text{unaffected}}$$

- Example: multiply accounts for 80s/100s
 - How much must multiply be improved to get 5× the overall performance?

$$20 = \frac{80}{n} + 20 \quad \quad \quad \blacksquare \text{ Can't be done!}$$

- Corollary: make the common case fast

MIPS measures integer performance

Engineering Applications employ floating point

- So the appropriate metric is Mega-flops (MFLOPS)
- Floating point support varies even more
- trig and exponentiation functions may be available
- In other cases, only simple arithmetic is supported

MFLOPS can be very unreliable in predicting performance

It is best to use execution time to assess performance