

Branches disrupt the flow of instructions in the pipeline

- The branch target address is computed in stage 3

- The branch condition is evaluated in stage 3

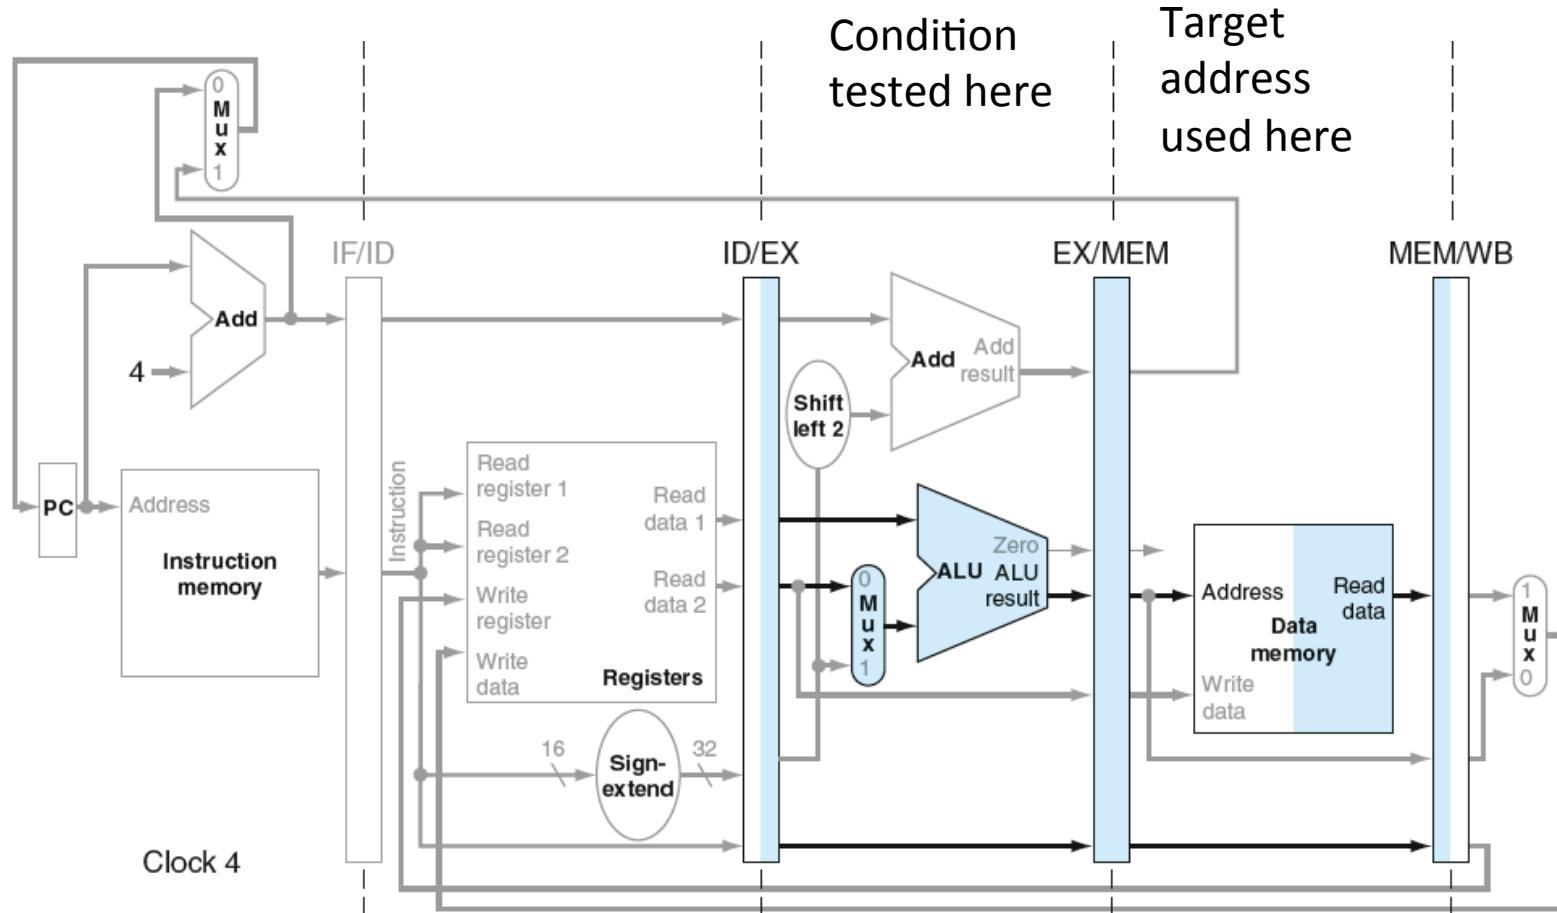
- The branch takes effect in stage 4

Instructions behind a taken branch must not complete

- Stages 1 through 3 must be flushed

- This creates 3 bubbles (a 3-cycle penalty)

The effect is called a control hazard



If condition is true, the PC is loaded with the target address when beq is in stage 4

sub \$11,\$6,\$5

beq \$11,\$0,skip

add \$4,\$7,\$3

sw \$9,4(\$7)

add \$9,\$2,\$9

.

.

.

skip: or \$8,\$4,\$0

If \$11 = 0, the 3 instructions that follow beq should not complete.
The next instruction to execute should be or \$8,\$4,\$0 instruction

Cycle	IF	ID	EX	MEM	WB
1	sub \$11,\$6,\$5				
2	beq \$11,\$0,skip	sub \$11,\$6,\$5			
3	add \$4,\$7,\$3	beq \$11,\$0,skip	sub \$11,\$6,\$5		
4	sw \$9,4(\$7)	add \$4,\$7,\$3	beq \$11,\$0,skip	sub \$11,\$6,\$5	
5	add \$9,\$2,\$9	sw \$9,4(\$7)	add \$4,\$7,\$3	beq \$11,\$0,skip	sub \$11,\$6,\$5
6	or \$8,\$4,\$0				beq \$11,\$0,skip

Flushing adds 3 extra clock cycles in this case

This is known as the branch penalty

1. Delay fetching instructions until branch behavior is known

Still causes 3-cycle penalty (outcome is known in stage 4)
2. Employ delayed branches

Fill delay slots with instructions needing to execute in any case
Compiler or programmer fills delay slots
Use NOPs if useful instructions cannot be identified
3. Evaluate the branch condition early

Requires computing target address in stage 2
Requires comparator in stage 2
4. Predict the behavior of the branch instruction

Requires recording previous branch behavior
Flushing is still required if prediction is wrong



The behavior of an unconditional branch is known

- The branch is decoded in stage 2

- The instruction following the branch is in stage 1

- This instruction could be flushed

Allowing the instruction to execute avoids a bubble

- The instruction is said to occupy the “branch delay slot”

- Compilers can move an instruction into the delay slot

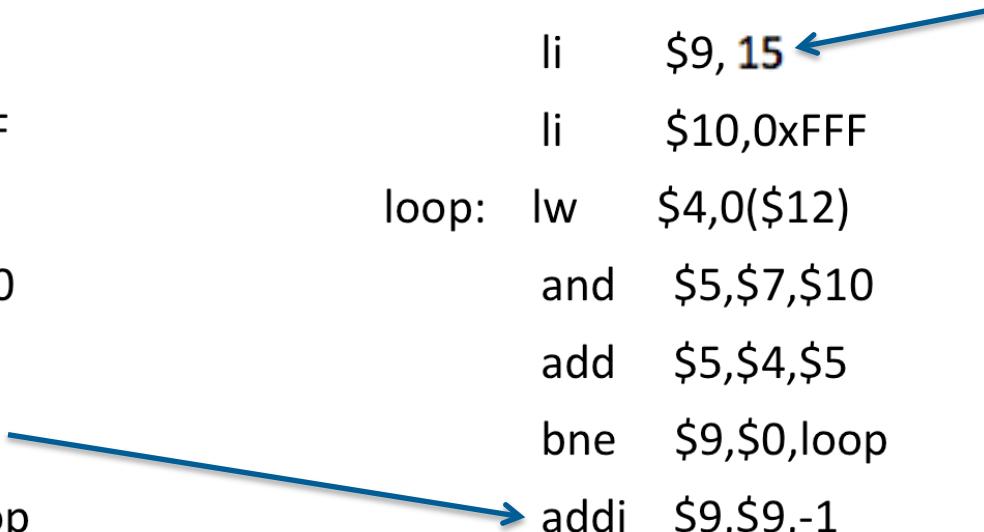
- It is easier to fill a single slot than to fill 3

li	\$9,16	li	\$9,16
li	\$10,0xFFFF	li	\$10,0xFFFF
loop:	lw \$4,0(\$12)	loop:	lw \$4,0(\$12)
	and \$5,\$7,\$10		and \$5,\$7,\$10
	add \$5,\$4,\$5		add \$5,\$4,\$5
	addi \$9,\$9,-1	bne \$9,\$0,loop	
	bne \$9,\$0,loop	addi \$9,\$9,-1	
	nop		



Branch delay slot instruction executes whether branch is taken or not
A cycle is saved by filling the slot with a useful instruction

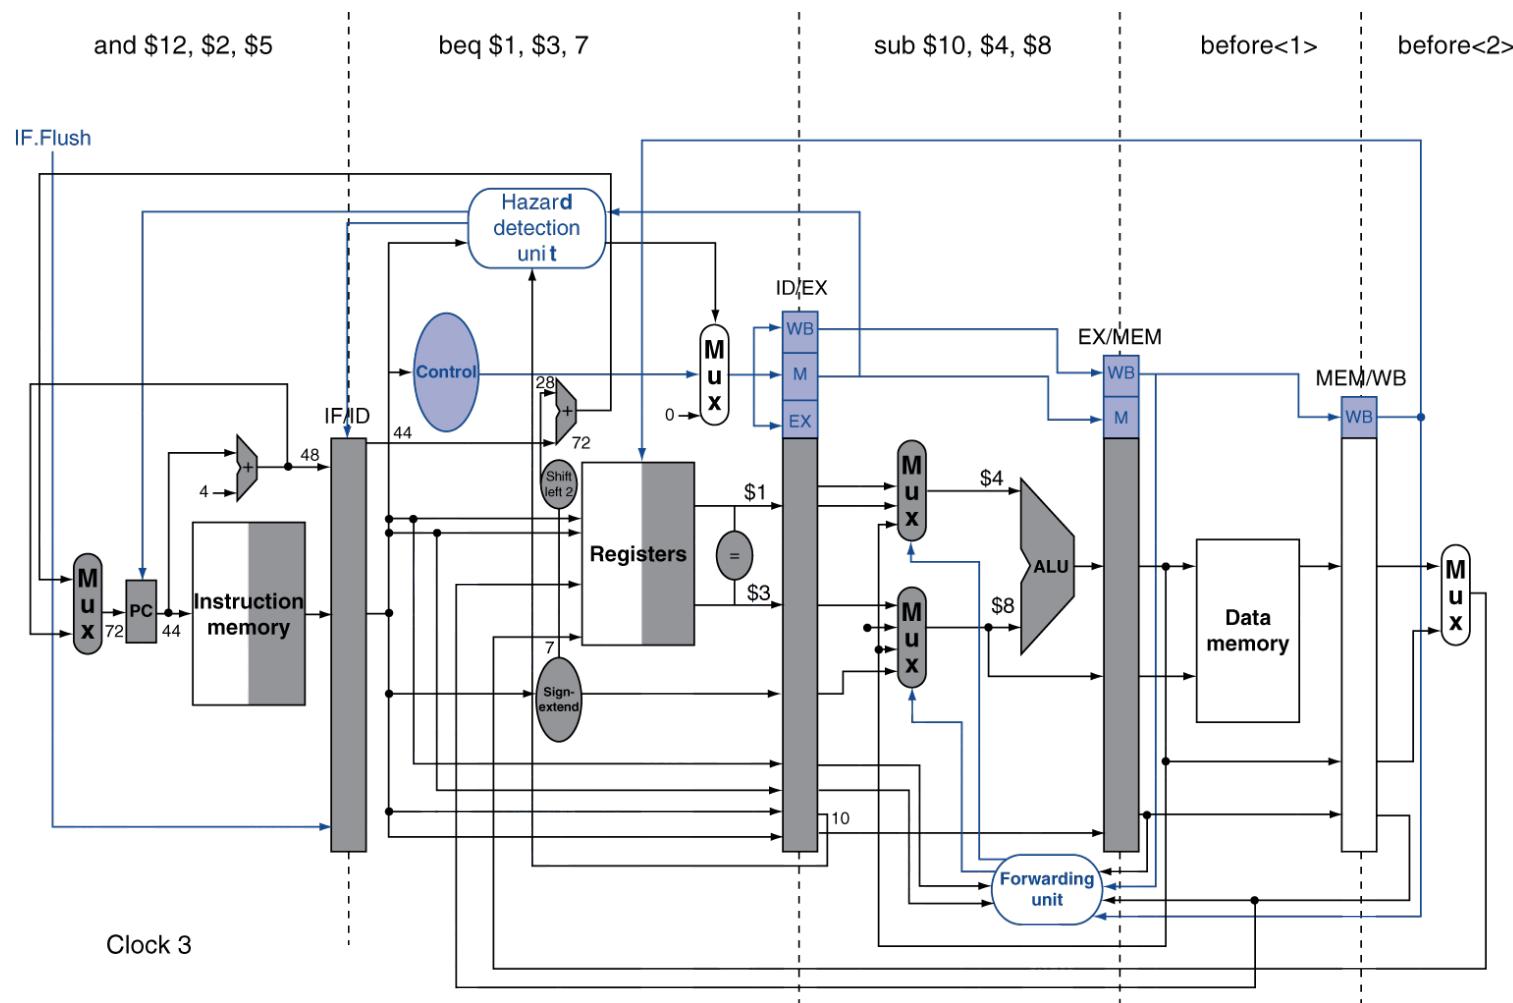
li	\$9,16	li	\$9, 15
li	\$10,0xFFFF	li	\$10,0xFFFF
loop:	lw \$4,0(\$12)	loop:	lw \$4,0(\$12)
	and \$5,\$7,\$10		and \$5,\$7,\$10
	add \$5,\$4,\$5		add \$5,\$4,\$5
	addi \$9,\$9,-1	bne \$9,\$0,loop	
	bne \$9,\$0,loop		addi \$9,\$9,-1
	nop		



The initial value in the loop control register has been set to 15 so that the number of loop iterations will be the same as for the original code.

- Move hardware to determine outcome to ID stage
 - Target address adder
 - Register comparator
- Example: branch taken

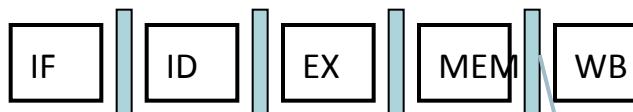
```
36: sub $10, $4, $8
40: beq $1, $3, 7
44: and $12, $2, $5
48: or $13, $2, $6
52: add $14, $4, $2
56: slt $15, $6, $7
...
72: lw $4, 50($7)
```



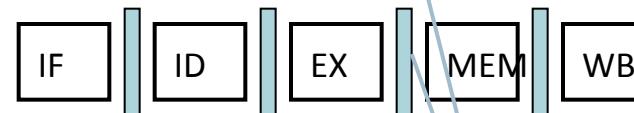
Delay slot is still needed since condition is tested in stage 2

- If a comparison register is a destination of 2nd or 3rd preceding ALU instruction

add \$1, \$2, \$3

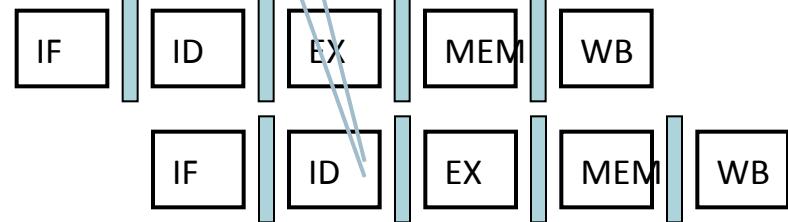


add \$4, \$5, \$6



...

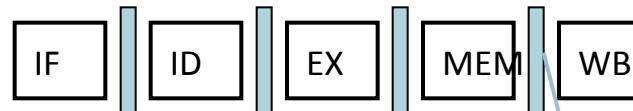
beq \$1, \$4, target



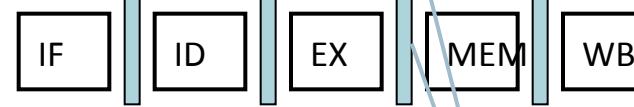
- Can resolve using forwarding

- If a comparison register is a destination of preceding ALU instruction or 2nd preceding load instruction
 - Need 1 stall cycle

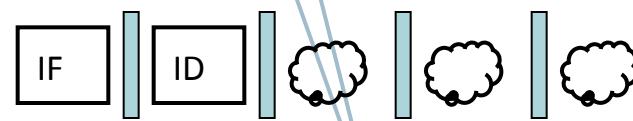
lw \$1, addr



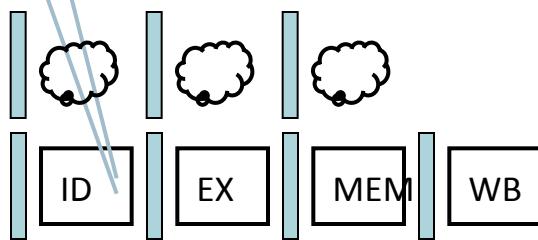
add \$4, \$5, \$6



beq stalled

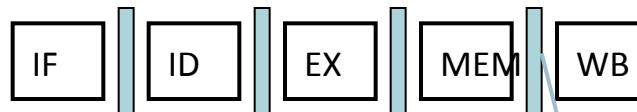


beq \$1, \$4, target



- If a comparison register is a destination of immediately preceding load instruction
 - Need 2 stall cycles

lw \$1, addr



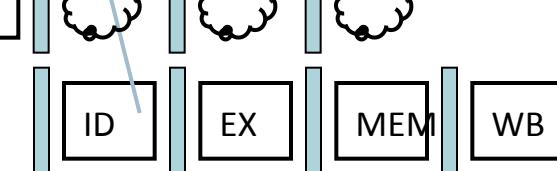
beq stalled



beq stalled



beq \$1, \$0, target



Branches have more impact on deeper and superscalar pipelines

More stages may have to be flushed

Superscalar pipelines process multiple instructions per stage

Branch prediction will be examined as another technique

Studies have shown:

20% to 30% of program instructions involve branching

About 65% of branches are taken

Instructions along the predicted path are speculative
the work done must be undone if the prediction is wrong

There are two options for branch prediction:

static prediction & dynamic prediction

Static prediction can be done by the compiler

Dynamic prediction requires extra hardware

Static (fixed) Prediction

The prediction is always the same (taken or not taken)

Forward branches may be predicted not taken

Backwards branches may be predicted taken
e.g. at end of loops

Once actual behavior is known, work may have to be undone

Dynamic Branch Prediction Increases accuracy of prediction

Previous history of branch behavior is recorded

Shows if branch was taken or not when last encountered

Branch prediction buffer

Branch History table (BHT)

Decode History table (DHT)

These tables are high-speed buffers (caches)

address of branch instruction is used to access them

BHT is accessed as soon as PC is updated

- Hits only occur for branch instructions already in the table

- Hits provide the predicted target address

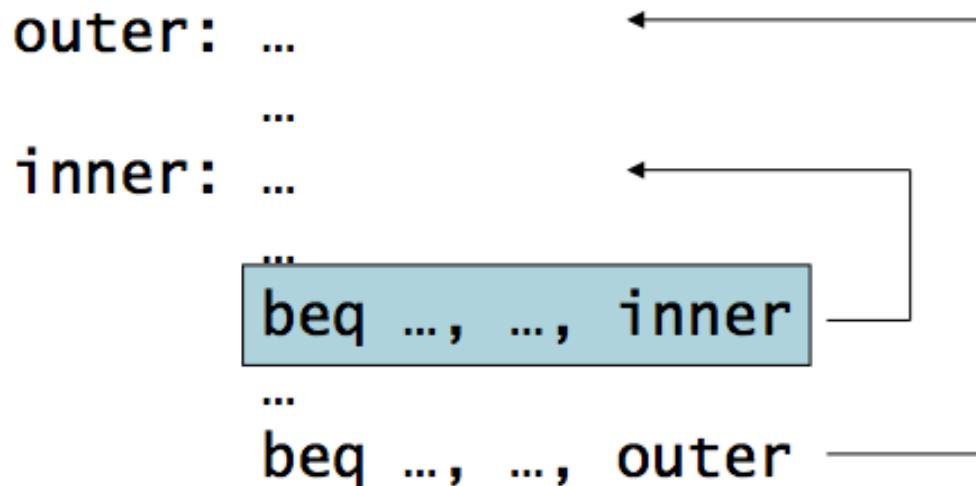
- Non-branch instructions cause misses

- Can be accessed before branch enters the pipeline

DHT is only accessed after branch is decoded

- It is not checked for non-branch instructions

- hits provide the predicted target instruction



Let's base prediction on a recorded bit (=0 if not taken, =1 if taken)

Assume 9 iterations of inner loop, 1st prediction & 9th are wrong

For each iteration of outer loop, inner beq is mispredicted twice

A 2-bit predictor would be better

2-bit predictor	Meaning
00	Strongly not taken
01	Weakly not taken
10	Weakly taken
11	Strongly taken

On 1st encounter bits change from 00 to 01 (for previous example)

On 2nd encounter bits change from 01 to 10

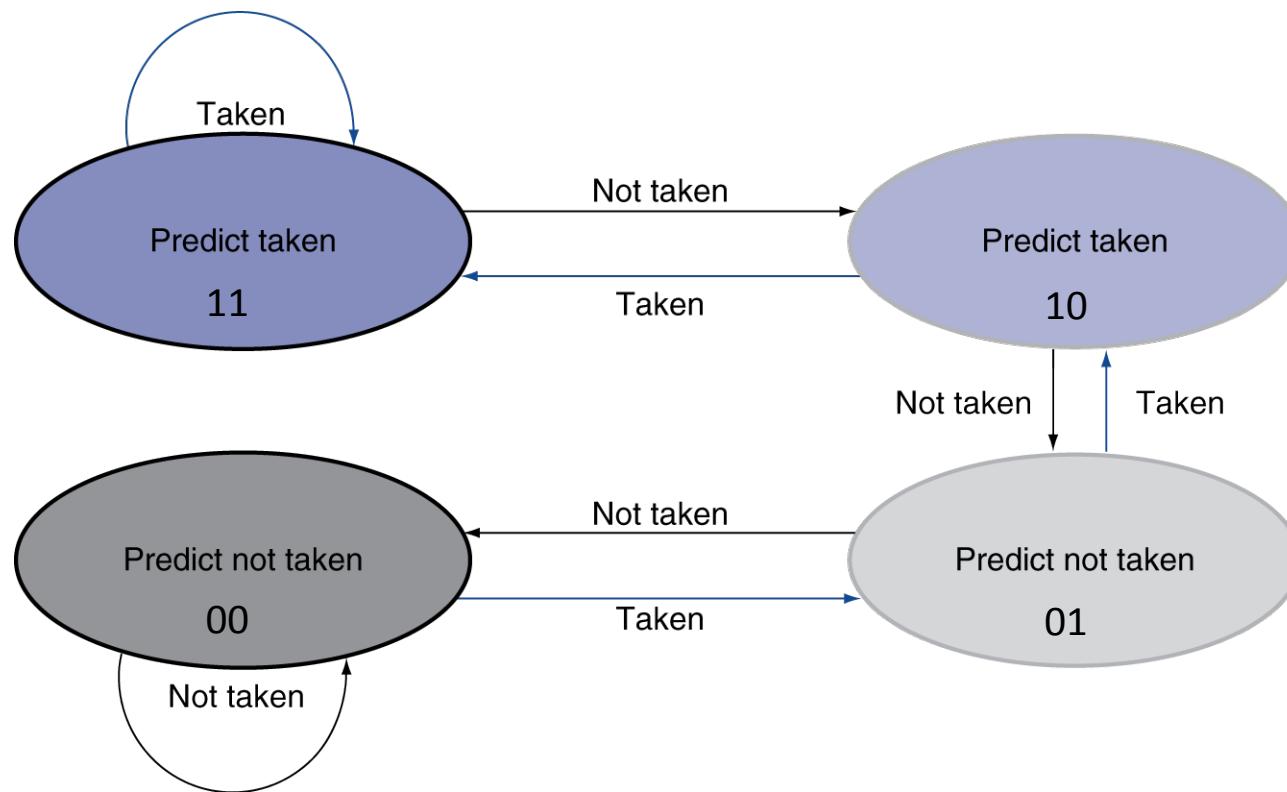
On 3rd encounter bits change from 10 to 11

On 9th encounter bits change from 11 to 10

Inner prediction is wrong three times for 1st iteration of outer loop

But wrong only once for each remaining outer loop iteration

- Only change prediction on two successive mispredictions



Exceptions affect the pipeline like a function call
control is diverted to the exception handler
interrupts are a particular type of exception

Exceptions are triggered by unexpected events
detecting invalid instructions during decode
arithmetic overflow
memory errors
syscall
interrupts from external I/O controllers

System Coprocessor (CP0) manages MIPS exceptions

EPC is the exception program counter (CP0 \$14)

Holds address of offending (or interrupted) instruction

Cause register (CP0 \$13) indicates problem

Control is transferred to handler at 0x800000018

handler takes the appropriate action

resumes program using address in EPC

terminates program if problem cannot be resolved

Transfer to exception handler causes a control hazard
program instructions are flushed from pipeline
this creates pipeline bubbles

The effect on pipeline is similar to a mispredicted branch

- Exception on add in

```
40      sub    $11,   $2,   $4
44      and    $12,   $2,   $5
48      or     $13,   $2,   $6
4C      add    $1,    $2,   $1
50      s1t    $15,   $6,   $7
54      lw     $16,  50($7)
```

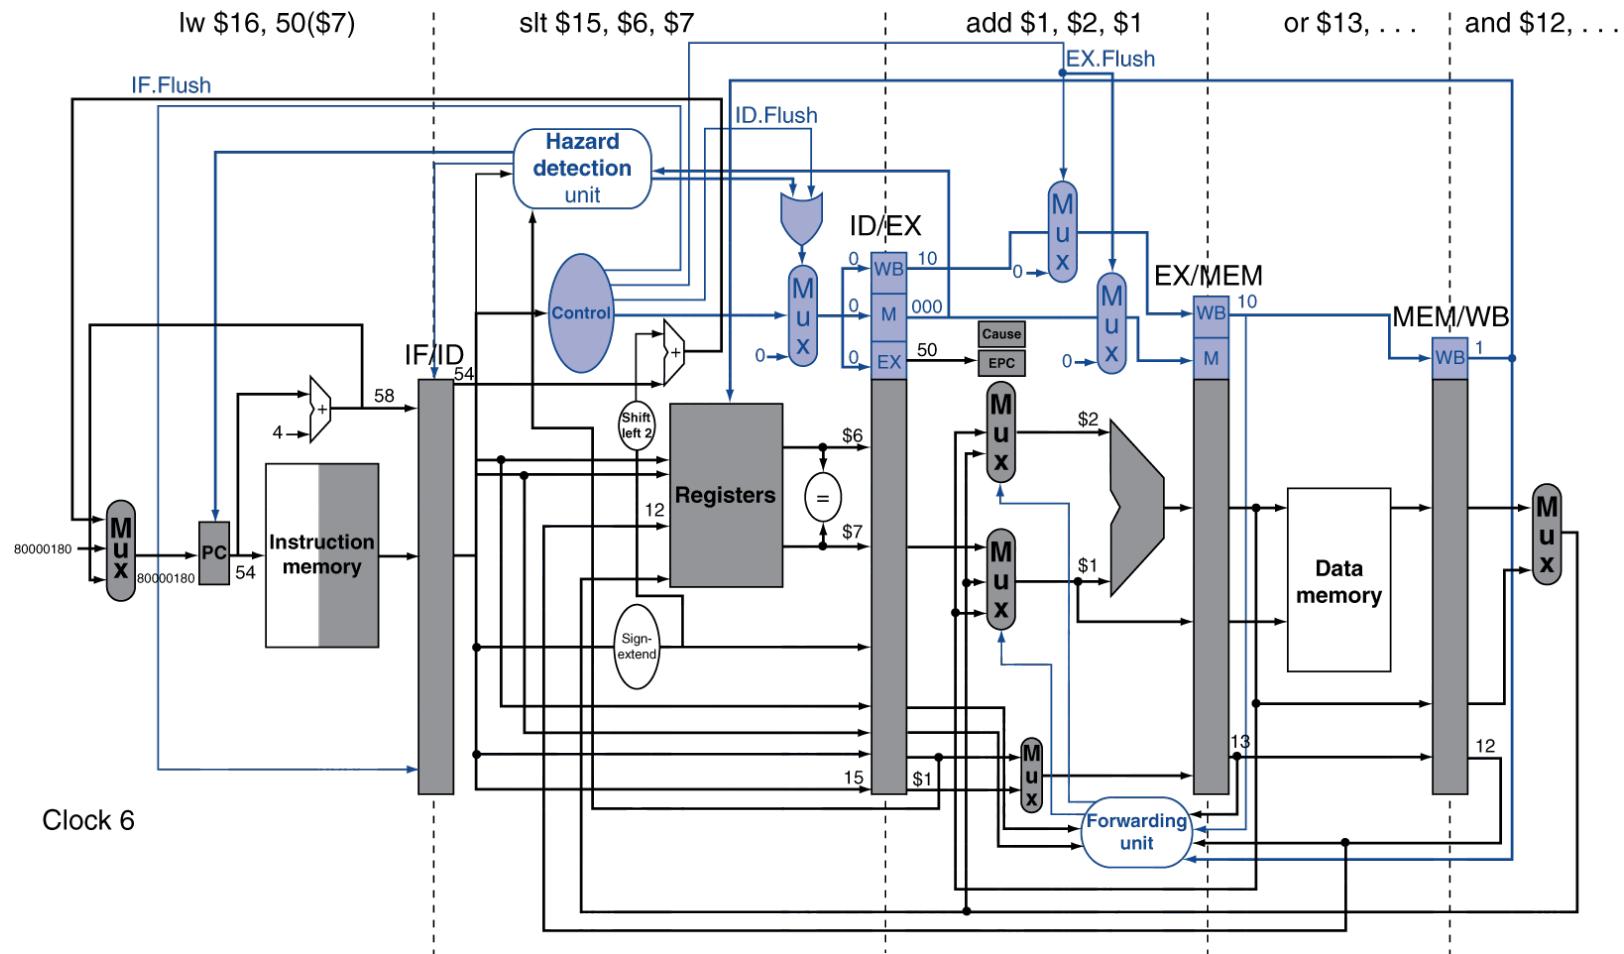
...

- Handler

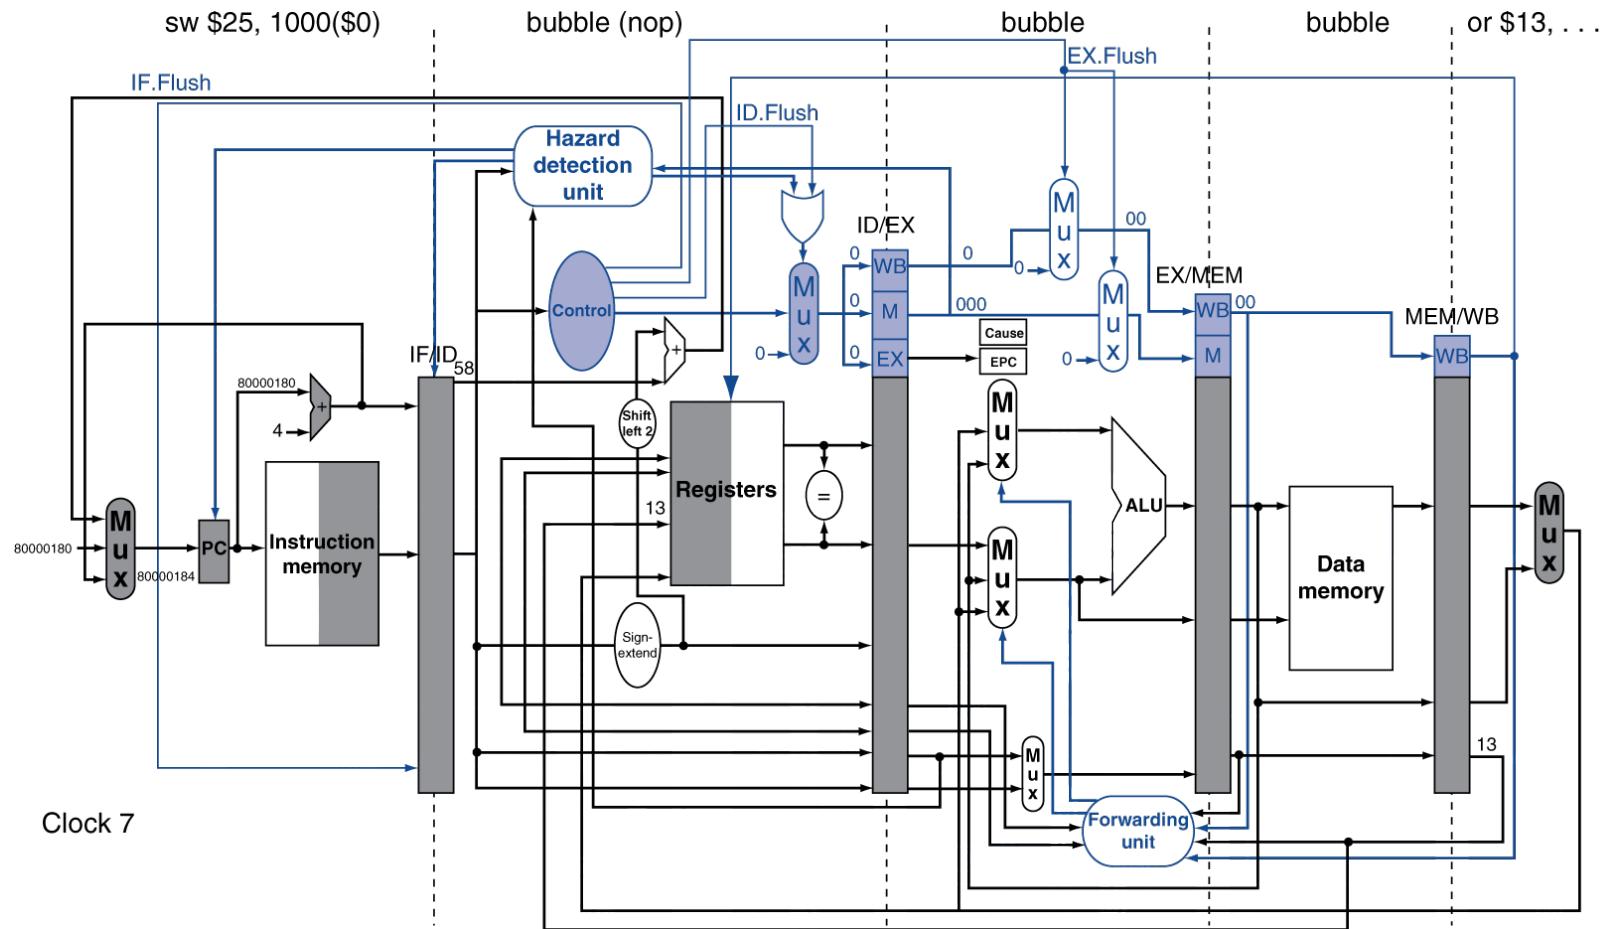
```
80000180      sw     $25,  1000($0)
80000184      sw     $26,  1004($0)
```

...

Exception Example



Flush control signal inserts bubbles



Flush control signal inserts bubbles

Some systems support precise exceptions
no register writes after the offending instruction
leaves the system in a consistent state

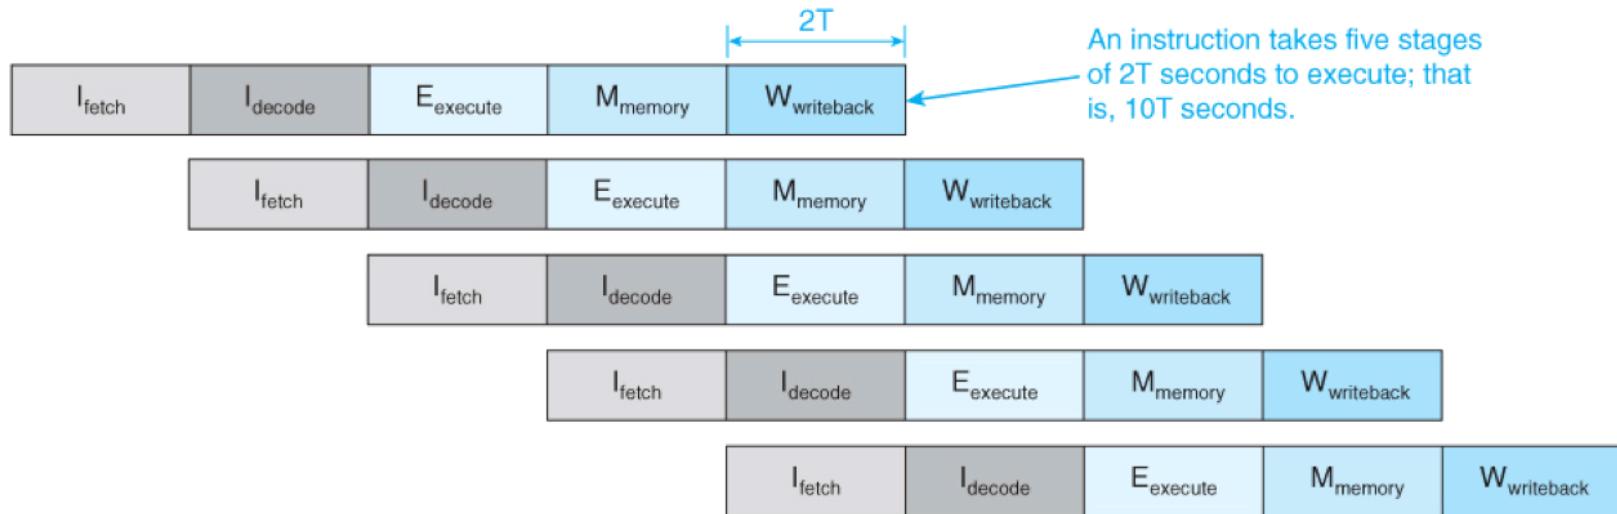
Imprecise exceptions
instruction causing exception is identified
but succeeding instructions may complete
allows the system to be in an inconsistent state

Superpipelining allows more instructions to be overlapped

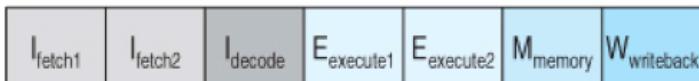
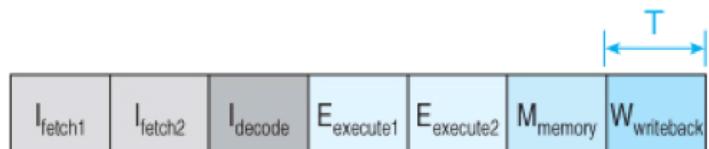
Operations may not all require a complete clock cycle
register reads or writes
check for hits in cache
decoding an opcode

The pipeline can run faster than the external clock rate

Superpipelines have some stages subdivided
a different instruction is in each phase with a stage



Conventional pipeline with cycle time = $2T$



The fetch stage has been divided into two pipelined stages.



An instruction takes five stages of $2T$, T , $2T$, T , T seconds, that is, $7T$ seconds.

The writeback stage requires only one cycle.



$$\text{Speedup} = 10/7 = 1.43$$

Superpipeline with cycle time = T (twice the rate)

Superpipelines use finer granularity
instruction throughput increases
the cost is a higher clock rate

Other parts of the datapath may require a slower rate
precludes running the entire system at a higher clock rate

Hazards & mispredictions have a greater impact
more stages have to be flushed

More stages cause more interstage delays
more pipeline registers

Superpipelines overlap more instructions

Superpipelining provides a modest performance increase

Superscalar provides a greater benefit

Superscalar systems can include superpipelining

Pipelining's goal is a throughput of 1 instruction per cycle
data hazards & control hazards make this difficult

Scalar systems handle 1 instruction per stage
each stage requires one clock cycle
only one instruction is started each cycle
all instructions go through all stages
some instructions may take longer than without the pipeline
however the instruction throughput is increased

Superscalar systems aim for 2 or more instructions per cycle
extra hardware executes multiple instruction per cycle
multiple pipelines can operate in parallel
or each stage can process multiple instructions at one time

Multiple instructions are fetched at one time
requires a wider CPU-to-memory bus

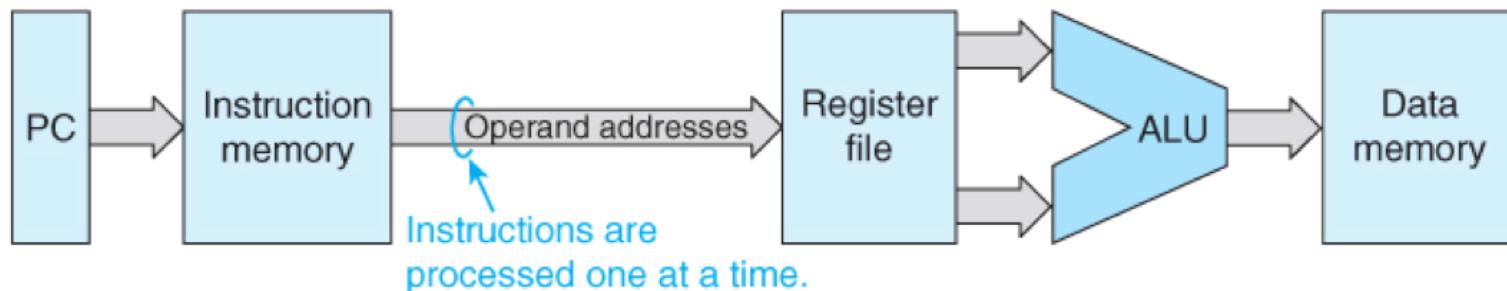
Multiple instructions are decoded together
dependencies between the instructions are detected
independent instructions are issued to their execute units

Multiple execute units process instructions at the same time
multiple ALUs for integer operations
multiple floating point units
load/store unit (to compute memory addresses)
branch unit to analyze branch conditions & make predictions

Resource Hazards are possible with superscalar systems
the required unit may not be available for next instruction
required registers may already be in use

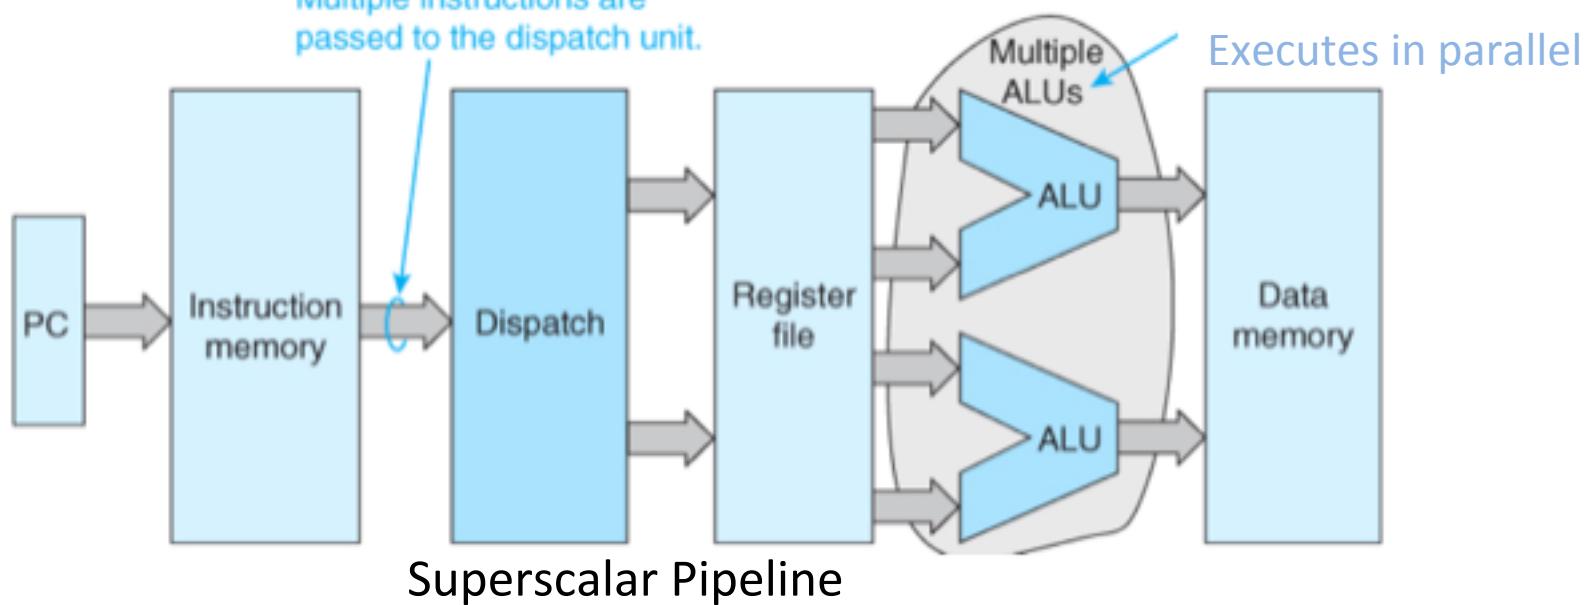
Out-of-order execution may be needed for good performance
instructions may start in a different order than in program
later instructions may complete before earlier ones

Scalar Pipeline



Multiple instructions are passed to the dispatch unit.

Executes in parallel



An m-way superscalar system has m parallel pipelines
m-fold increases in performance are seldom achieved

The Pentium P5 had 2 integer pipelines
the 2nd pipeline could only be used about 30% of the time

The original Pentium used superscalar operation
two 5-stage integer pipelines U and V
U pipe included a shifter not in the V pipe
a 6-stage floating point unit

Compilers generate machine instructions in program order

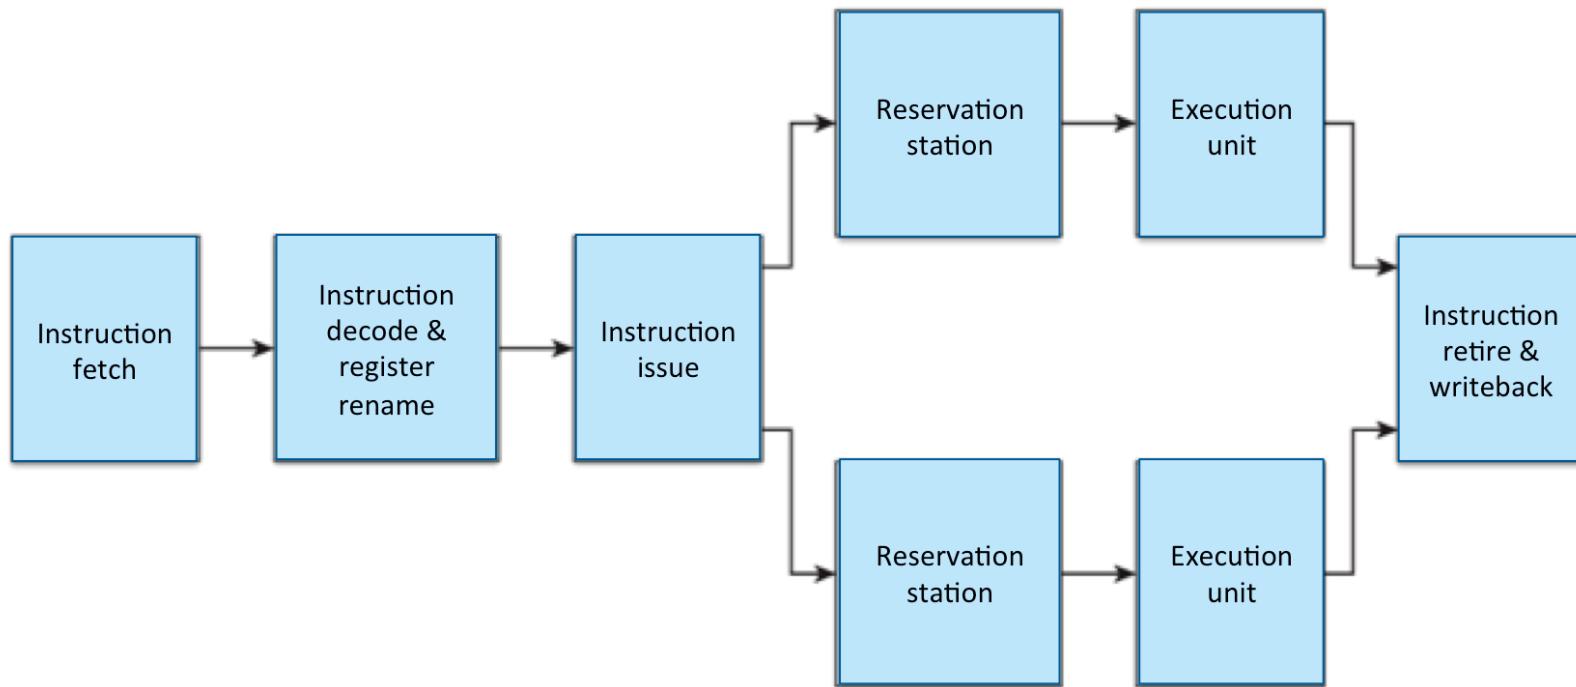
Superscalar processors *dispatch* multiple instructions in parallel
operands are obtained from multiported register files
Instructions that have operands can execute in parallel
The execute units needed must be available

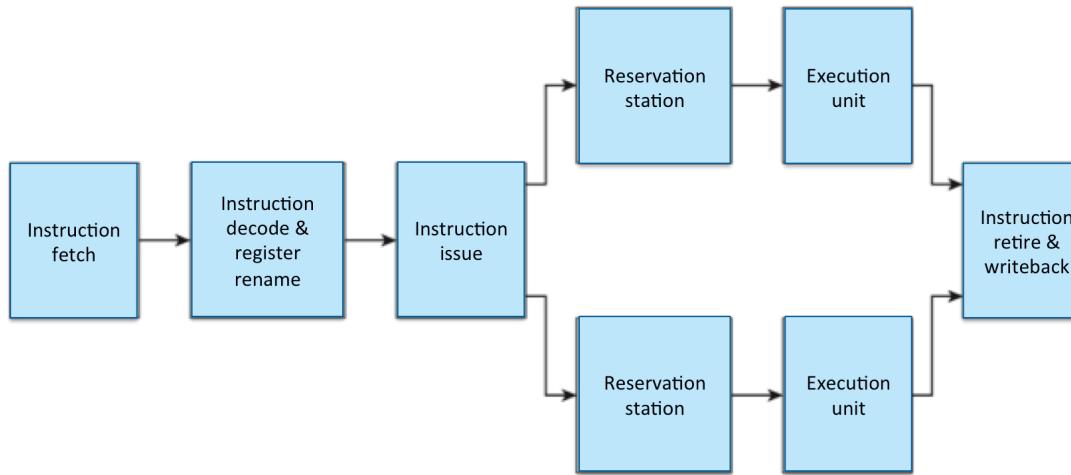
The program semantics (meaning must be preserved)

	Sequence 1	Sequence 2
add	\$11,\$12,\$13	\$11,\$12,\$13
add	\$14,\$11,\$13	\$15,\$16,\$17
add	\$15,\$16,\$17	\$14,\$11,\$13

Different instruction order but same meaning or net effect

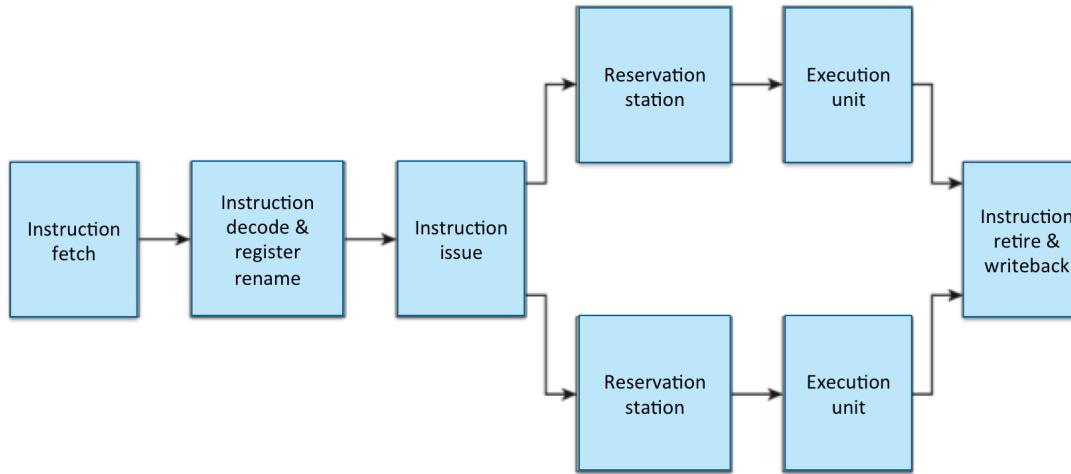
A more realistic view of a superscalar system is:





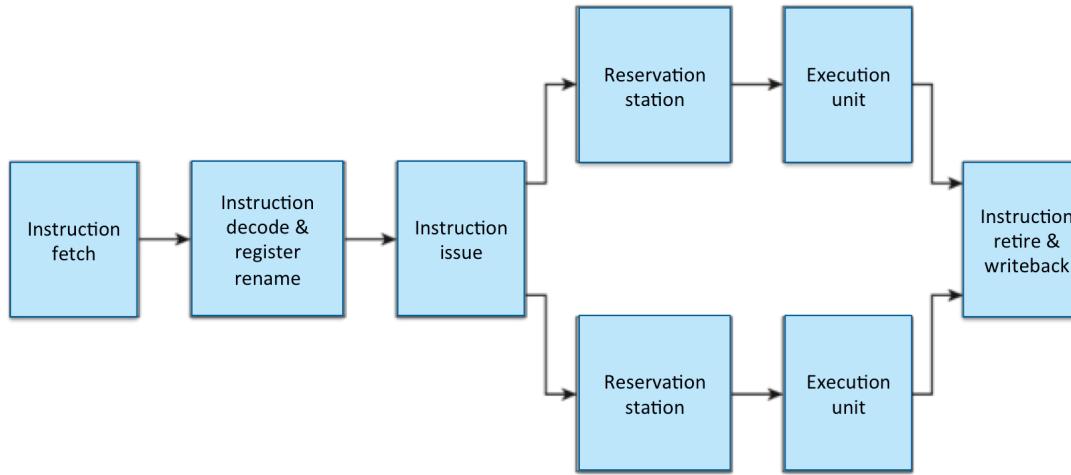
Instruction fetch
obtains instructions from cache or memory

Instruction Decode
interprets opcodes
substitutes temporary registers to avoid unnecessary stalls



Instruction issue

ensures as many instructions as possible execute in parallel
sends instructions to reservation stations (*issues* them)
Instructions are *dispatched* from the reservation stations
once the required input operands are available



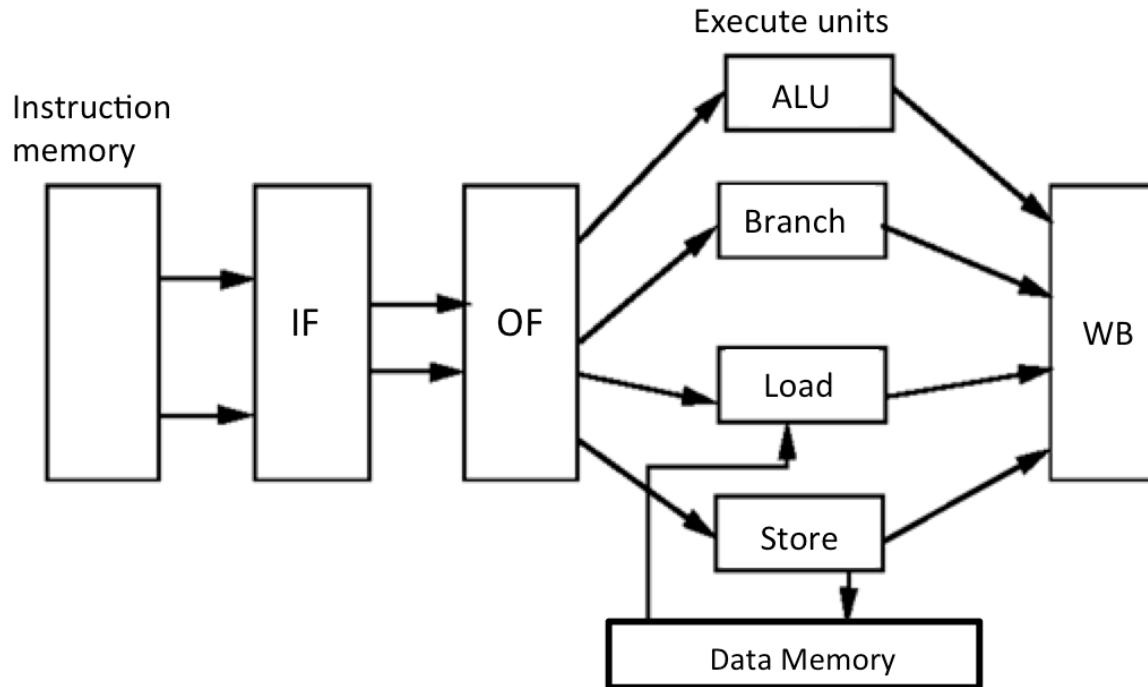
Reservation stations

Serve as front end buffer to execution units

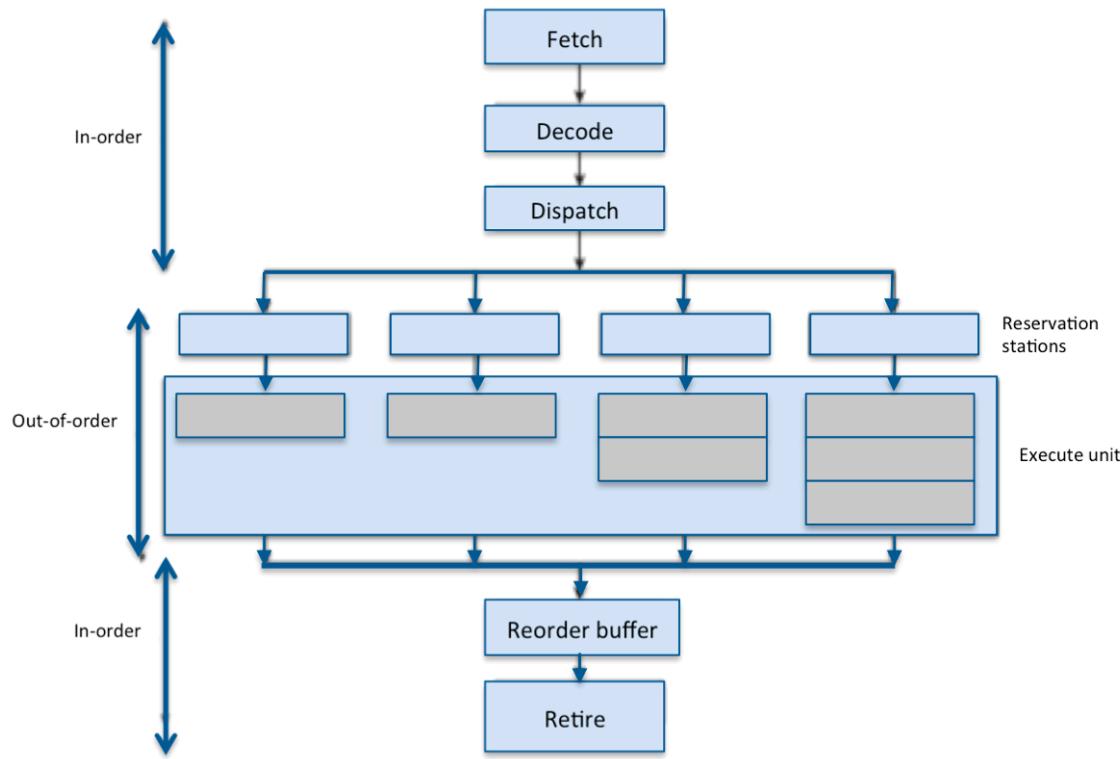
Hold instructions that use the corresponding execution unit

Instruction retire

Writes the results to the destination registers (*commits* them)
Tells reservation stations when resources are available



Even with in-order issue, instructions may complete out-of-order
e.g. An add or shift may be issued after a lw or floating point
instruction, but may finish executing first.



Reorder buffers hold completed instructions
ensure that the program meaning or outcome is preserved
instructions that complete out-of-order are retired in-order
retire means writing the result to the correct register

Possible issue/completion options:

in-order issue & in-order completion

 used in scalar pipeline

in-order issue & out-of-order completion

out-of-order issue & in-order completion

out-of-order issue & out-of-order completion

Changing execution order can cause other types of data dependencies

1. True data dependency (RAW) read after write

Sub \$3, \$2, \$4

Add \$5,\$3,\$4 # must read \$5 after it is written by sub

2. Antidependency (WAR) write after read

or \$2, \$4, \$3

sub \$4, \$5,\$6 # can't write \$4 until after the or reads \$4

The or may be stalled waiting on a resource or on \$3

3. Output dependency (WAW) write after write

sub \$3,\$2,\$4

add \$5,\$6,\$5

slt \$3,\$7,\$0 # must write \$3 after sub writes \$3

The use of registers can create apparent dependencies

1. or \$3, \$3,\$5
2. add \$4,\$3,\$2
3. add \$3, \$5, \$2
4. sub \$7,\$3,\$4

WAW exists between 3. and 1. (both write \$3)

WAR exists between 3. and 2. (3. overwrites \$3 input to 2.)

RAW exists between (1. & 2.), (2. & 4.), (3. & 4.)

The use of registers can create apparent dependencies

1. or \$3_a, \$3_a, \$5
2. add \$4, \$3_a, \$2
3. add \$3_b, \$5, \$2
4. sub \$7, \$3_b, \$4

3_a and 3_b are unrelated, so a different register can be used for 3_b

Eliminates false dependencies

Register renaming performs this dynamic substitution

Register renaming adapts to larger register sets

Applications use more registers without having to be recompiled

Recompilation is required without register renaming feature

Instruction-level Parallelism (IPL)
the ability to execute multiple instructions together

Pipelining provides IPL by overlapping instruction execution

Superscalar systems replicate hardware to provide IPL
multiple instructions can be sent through separate pipelines
multiple units can execute instructions simultaneously

Multiple issue means starting more than 1 instruction at a time

Static multiple issue is based on:

compiler deciding which instruction can execute together

Independent instructions are grouped into packets or bundles

This packaging occurs prior to execution (compile time)

Bundles or packets are assigned to issue slots

These are also called very long instruction words (VLIW)

Packets are issued within a single clock cycle

Packed instructions must map to separate resources

Instruction mix may be restricted

Dynamic multiple issue is based on:

Control unit deciding which instruction can execute together

Decisions are made by hardware during execution

Based on hazards (data, resources, etc.)

Compiler may assist by reordering machine instructions

Consider a 2-issue MIPS processor

Two instructions are fetched and executed together

Requires 64-bit CPU-to-memory bus

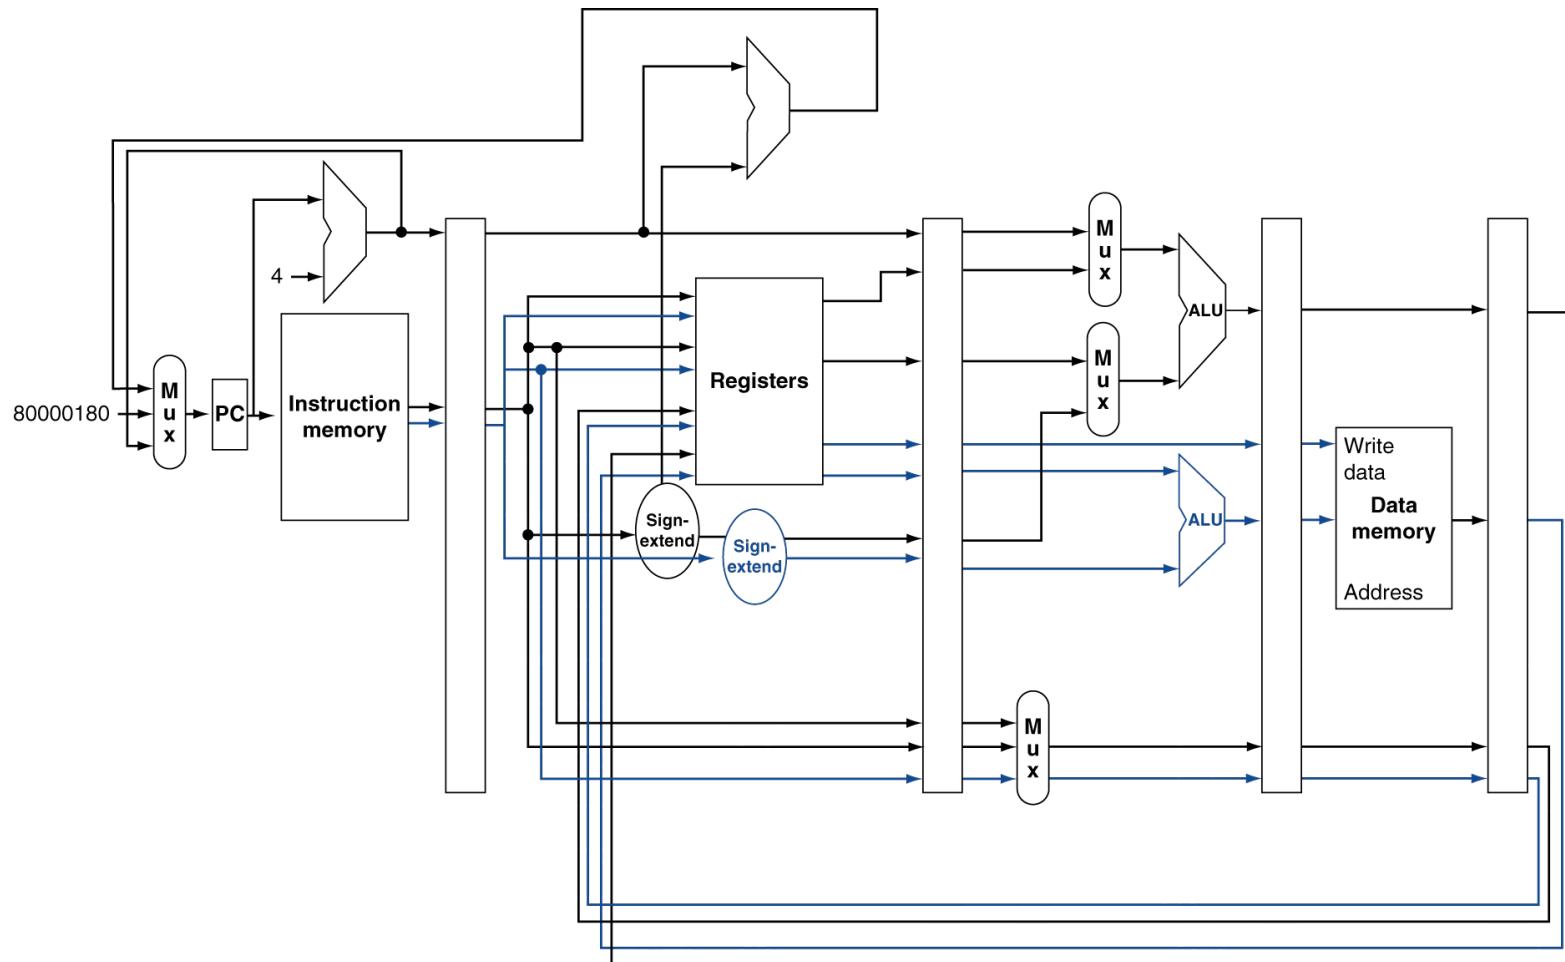
One can be an integer or branch instruction

The other can be a load or store instruction

Instructions are always issued in pairs

Nop instruction replaces an unused slot within pair

Required stalls holdup both instructions in a pair



Extra resources: updated register file, unit to compute memory address,
2nd sign extension unit

- More instructions executing in parallel
- EX data hazard
 - Forwarding avoided stalls with single-issue
 - Now can't use ALU result in load/store in same packet
 - add \$t0, \$s0, \$s1
 - load \$s2, 0(\$t0)
 - Split into two packets, effectively a stall
- Load-use hazard
 - Still one cycle use latency, but now two instructions
- More aggressive scheduling required

```
Loop: lw    $t0, 0($s1)      # $t0=array element
      addu $t0, $t0, $s2      # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi $s1, $s1,-4        # decrement pointer
      bne  $s1, $zero, Loop  # branch $s1!=0
```

	ALU/branch	Load/store	cycle
Loop:	nop	lw \$t0, 0(\$s1)	1
	addi \$s1, \$s1,-4	nop	2
	addu \$t0, \$t0, \$s2	nop	3
	bne \$s1, \$zero, Loop	sw \$t0, 4(\$s1)	4

- IPC = 5/4 = 1.25 (c.f. peak IPC = 2)



- Replicate loop body to expose more parallelism
 - Reduces loop-control overhead
- Use different registers per replication
 - Called “register renaming”
 - Avoid loop-carried “anti-dependencies”
 - Store followed by a load of the same register
 - Aka “name dependence”
 - Reuse of a register name

	ALU/branch	Load/store	cycle
Loop:	addi \$s1, \$s1,-16	lw \$t0, 0(\$s1)	1
	nop	lw \$t1, 12(\$s1)	2
	addu \$t0, \$t0, \$s2	lw \$t2, 8(\$s1)	3
	addu \$t1, \$t1, \$s2	lw \$t3, 4(\$s1)	4
	addu \$t2, \$t2, \$s2	sw \$t0, 16(\$s1)	5
	addu \$t3, \$t4, \$s2	sw \$t1, 12(\$s1)	6
	nop	sw \$t2, 8(\$s1)	7
	bne \$s1, \$zero, Loop	sw \$t3, 4(\$s1)	8

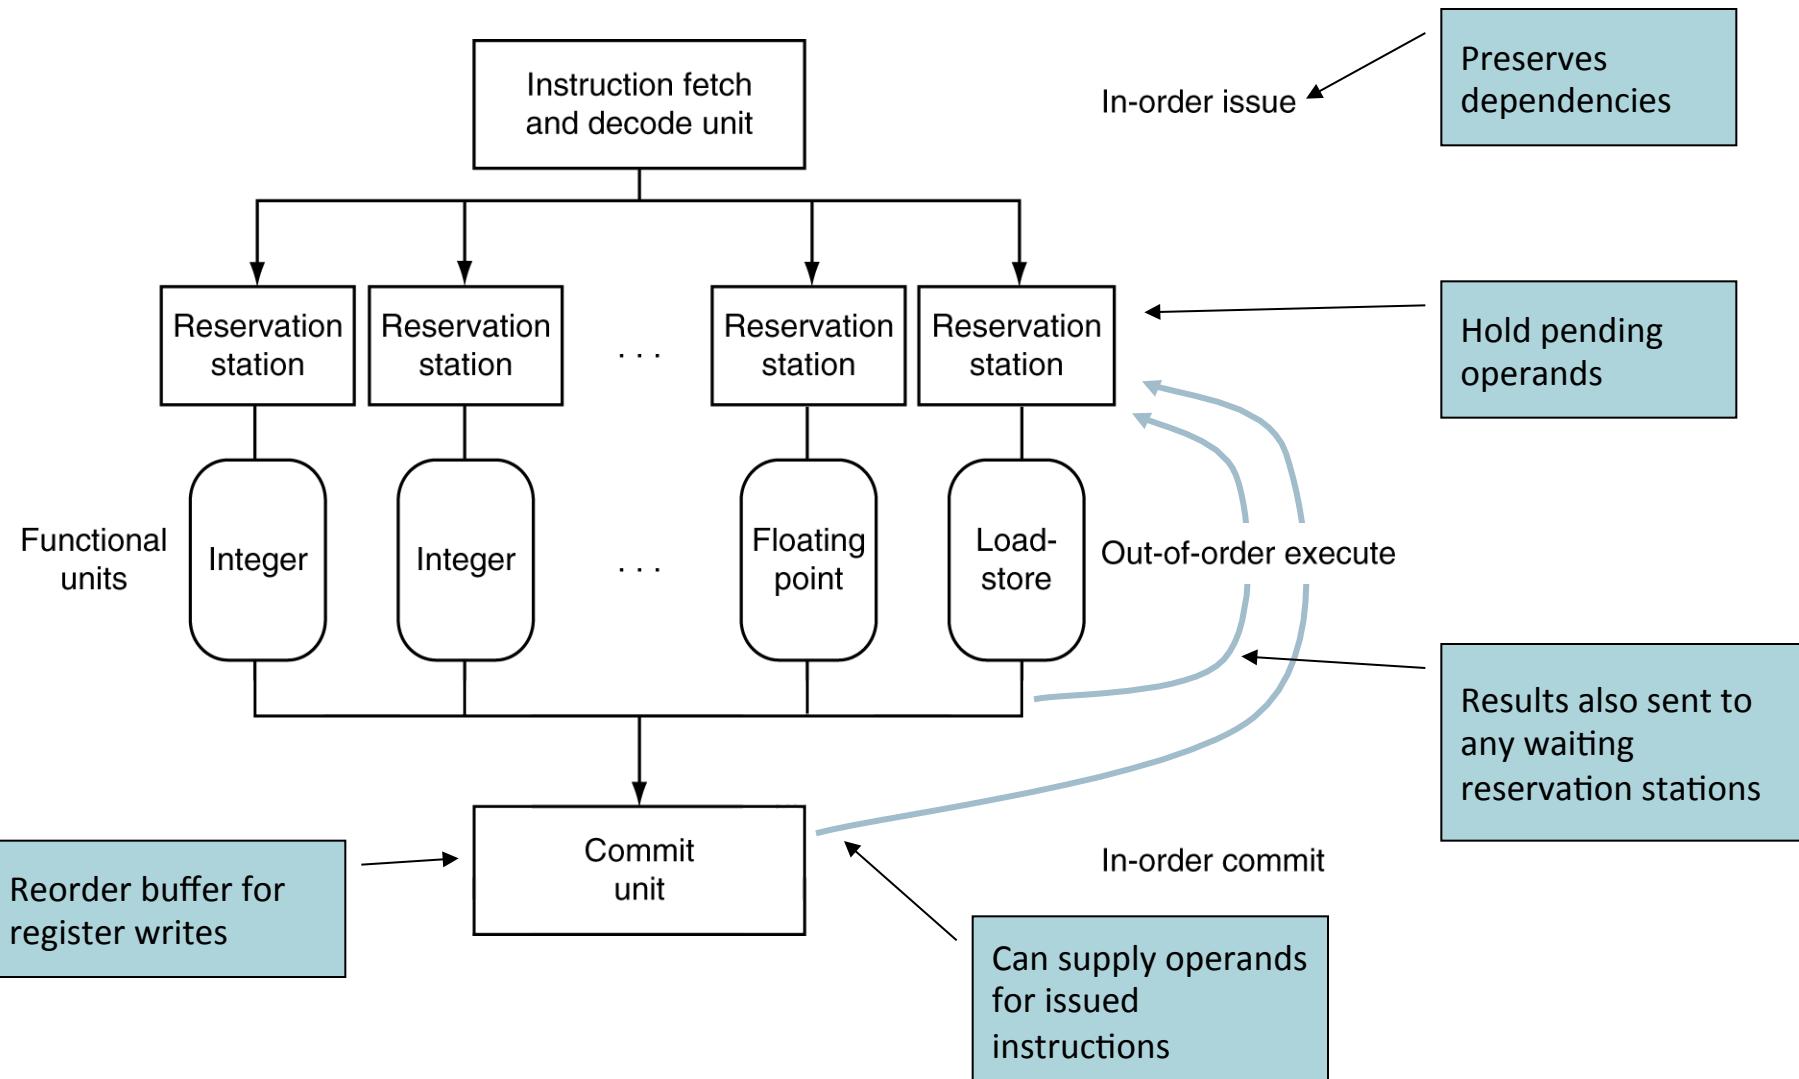
- $IPC = 14/8 = 1.75$
 - Closer to 2, but at cost of registers and code size

- “Superscalar” processors
- CPU selects instructions to execute each cycle
 - Avoiding structural and data hazards
- Avoids the need for compiler scheduling
 - Though compiler may still help
 - Code semantics ensured by the CPU

- Allow the CPU to execute instructions out of order to avoid stalls
 - But commit result to registers in order
- Example

```
lw      $t0, 20($s2)
addu   $t1, $t0, $t2
sub    $s4, $s4, $t3
slt   $t5, $s4, 20
```

- Can start sub while addu is waiting for lw



- Register renaming via reservation stations (RS) & ROB
 - ROB is reorder buffer
- On instruction issue to reservation station
 - copy available operand to reservation station
 - from register file or from ROB
 - Once done, register can be overwritten
 - When available, operands are provided to RS
 - Come from the function unit producing the result
 - Register update may not be required


```
Loop: lw    $t0, 0($s1)      # $t0=array element
      addu $t0, $t0, $s2      # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi $s1, $s1,-4        # decrement pointer
      bne  $s1, $zero, Loop  # branch $s1!=0
```

	ALU/branch	Load/store	cycle
Loop:	nop	lw \$t0, 0(\$s1)	1
	addi \$s1, \$s1,-4	nop	2
	addu \$t0, \$t0, \$s2	nop	3
	bne \$s1, \$zero, Loop	sw \$t0, 4(\$s1)	4

- IPC = 5/4 = 1.25 (c.f. peak IPC = 2)



- Replicate loop body to expose more parallelism
 - Reduces loop-control overhead
- Use different registers per replication
 - Called “register renaming”
 - Avoid loop-carried “anti-dependencies”
 - Store followed by a load of the same register
 - Aka “name dependence”
 - Reuse of a register name

Loop:

```
lw    $t0, 0($s1)
addu $t0, $t0, $s2
sw    $t0, 0($s1)
addi $s1, $s1,-4
lw    $t0, 0($s1)
addu $t0, $t0, $s2
sw    $t0, 0($s1)
addi $s1, $s1,-4
lw    $t0, 0($s1)
addu $t0, $t0, $s2
sw    $t0, 0($s1)
addi $s1, $s1,-4
lw    $t0, 0($s1)
addu $t0, $t0, $s2
sw    $t0, 0($s1)
addi $s1, $s1,-4
lw    $t0, 0($s1)
addu $t0, $t0, $s2
sw    $t0, 0($s1)
addi $s1, $s1,-4
bne  $s1, $zero, Loop
```

Loop:

```
addi $s1, $s1,-16
lw    $t0, 16($s1)
addu $t0, $t0, $s2
sw    $t0, 16($s1)
lw    $t1, 12($s1)
addu $t1, $t1, $s2
sw    $t1, 12($s1)
lw    $t2, 8($s1)
addu $t2, $t2, $s2
sw    $t2, 8($s1)
lw    $t3, 4($s1)
addu $t3, $t3, $s2
sw    $t3, 4($s1)
bne  $s1, $zero, Loop
```



	ALU/branch	Load/store	cycle
Loop:	addi \$s1, \$s1,-16	lw \$t0, 0(\$s1)	1
	nop	lw \$t1, 12(\$s1)	2
	addu \$t0, \$t0, \$s2	lw \$t2, 8(\$s1)	3
	addu \$t1, \$t1, \$s2	lw \$t3, 4(\$s1)	4
	addu \$t2, \$t2, \$s2	sw \$t0, 16(\$s1)	5
	addu \$t3, \$t4, \$s2	sw \$t1, 12(\$s1)	6
	nop	sw \$t2, 8(\$s1)	7
	bne \$s1, \$zero, Loop	sw \$t3, 4(\$s1)	8

(a 0 displacement & the original value in \$s1 are used in the first lw instruction)

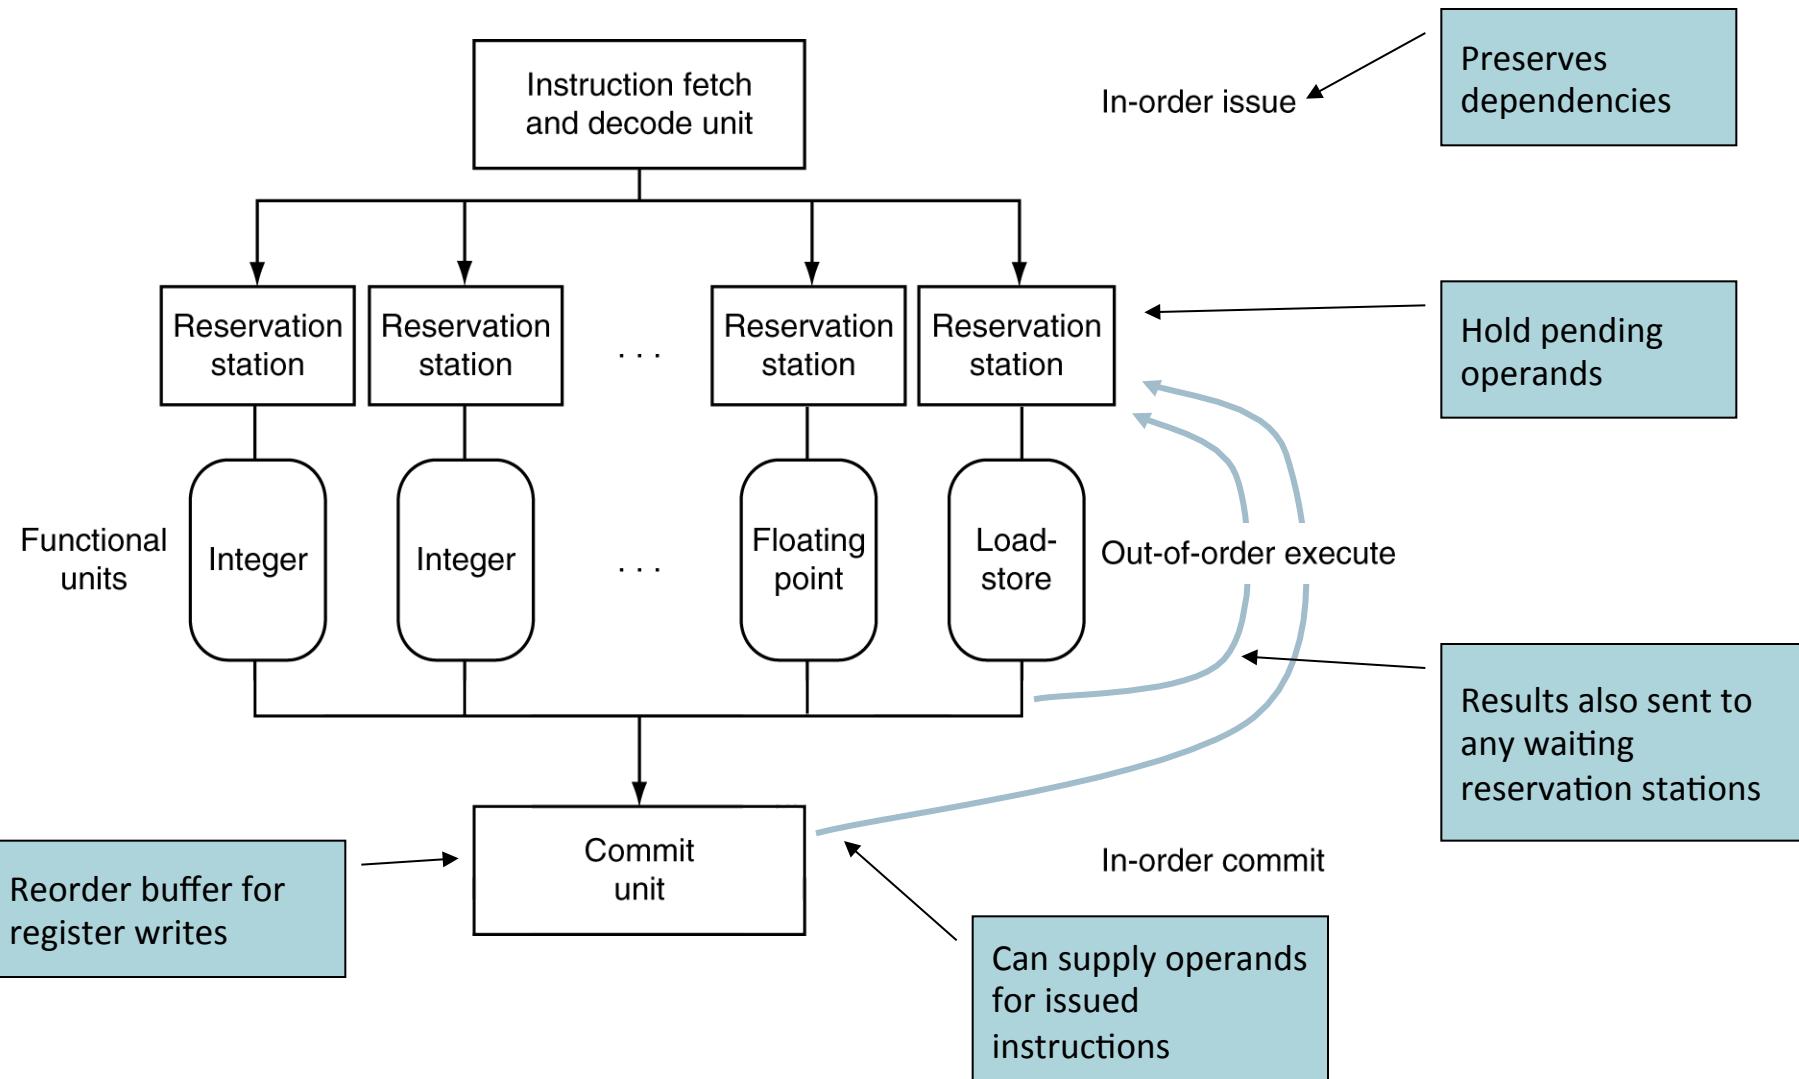
- $\text{IPC} = 14/8 = 1.75$
 - Closer to 2, but at cost of registers and code size

- “Superscalar” processors
- CPU selects instructions to execute each cycle
 - Avoiding structural and data hazards
- Avoids the need for compiler scheduling
 - Though compiler may still help
 - Code semantics ensured by the CPU

- Allow the CPU to execute instructions out of order to avoid stalls
 - But commit result to registers in order
- Example

```
lw      $t0, 20($s2)
addu   $t1, $t0, $t2
sub    $s4, $s4, $t3
slti   $t5, $s4, 20
```

- Can start sub while addu is waiting for lw



- Register renaming via reservation stations (RS) & ROB
 - ROB is reorder buffer
- On instruction issue to reservation station
 - copy available operand to reservation station
 - from register file or from ROB
 - Once done, register can be overwritten
 - When available, operands are provided to RS
 - Come from the function unit producing the result
 - Register update may not be required