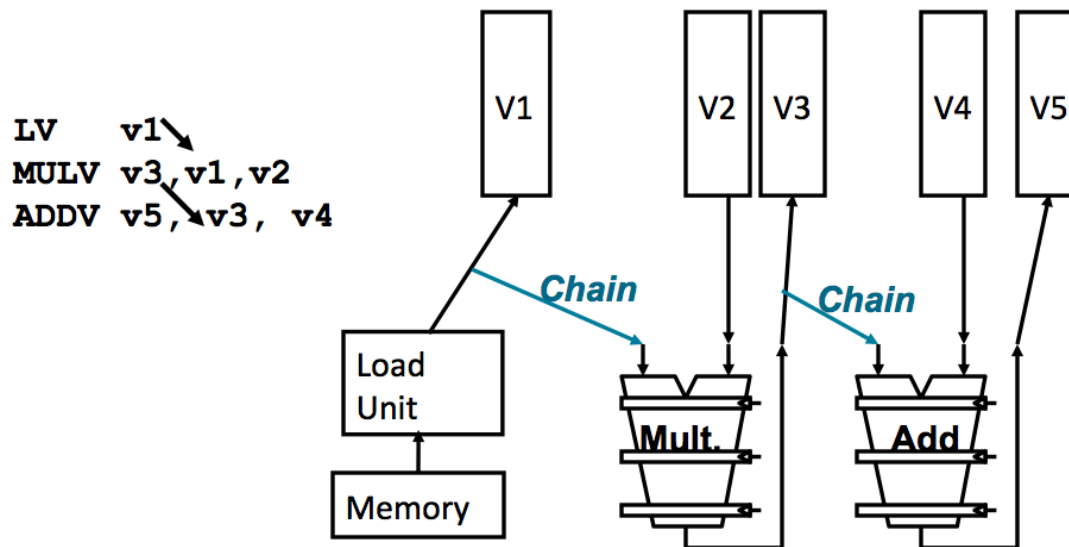


Input operands are used as soon as they arrive from memory  
Avoids stalling until the entire vector register is filled

Results can be sent from one functional unit directly to another  
A vector version of register bypassing and forwarding



Chaining overlaps the loading of operands into vector registers

The next operands are read in parallel with the use of the previous set

The vector length register (VLR) specifies number of reads

Vector elements do not have to occupy adjacent memory cells

Unlike with multi-media extensions (MMX, SSE and AVX)

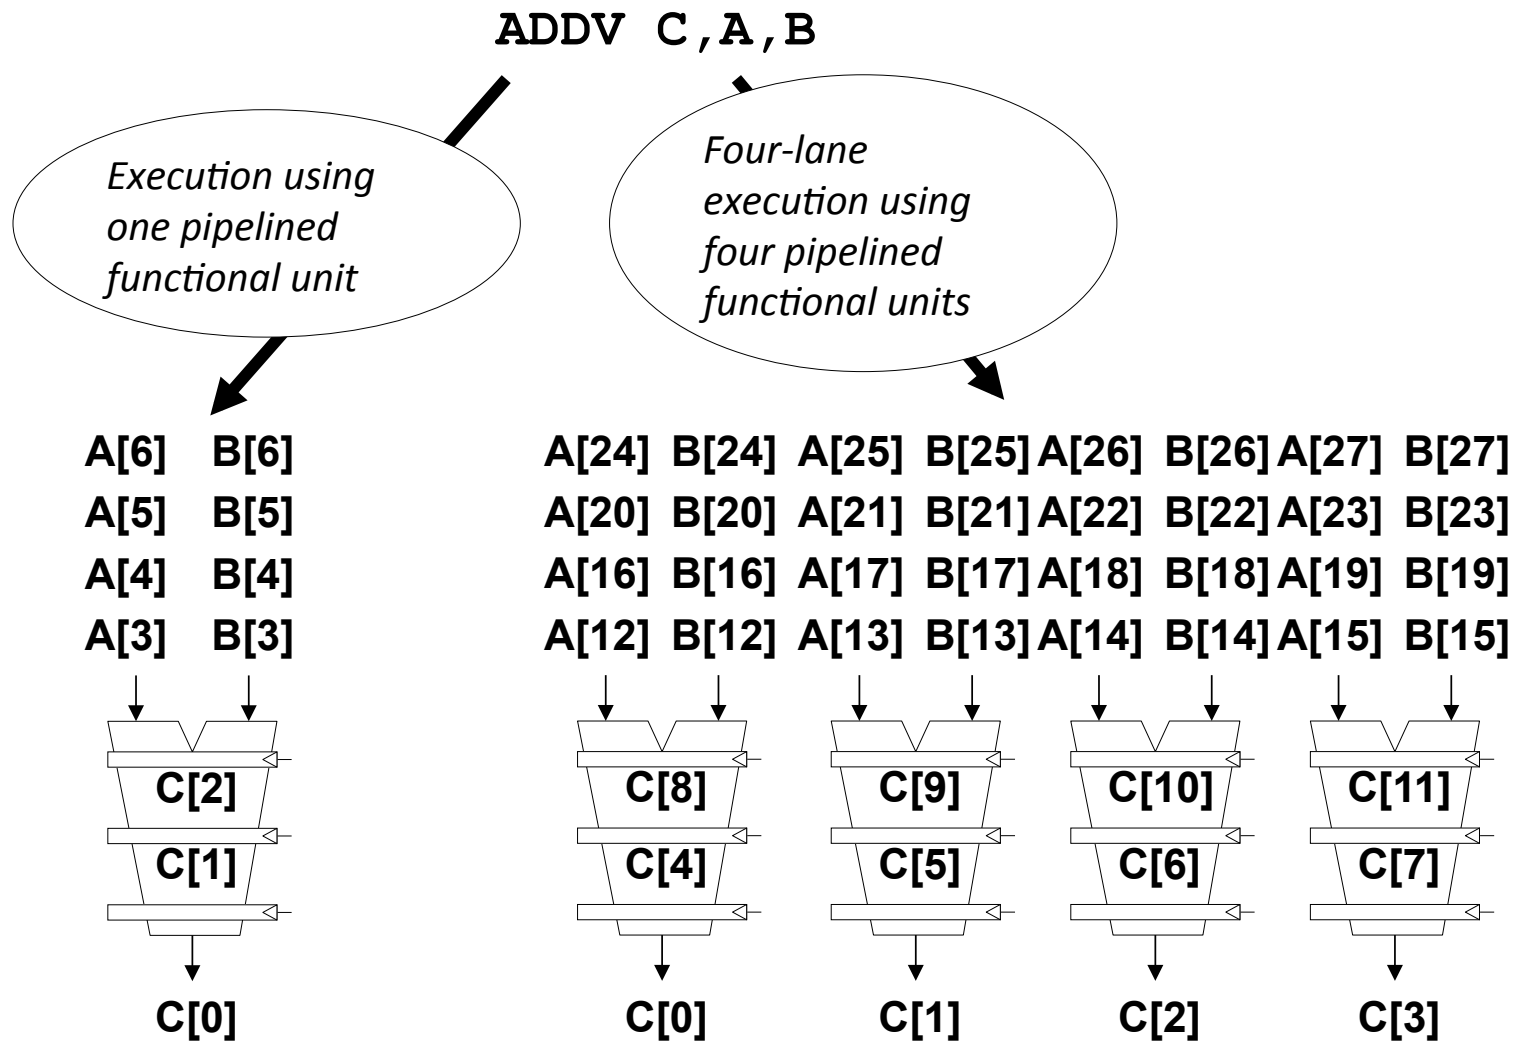
*Stride* = amount by which elements are separated in memory

Results can be chained back to memory as they are produced

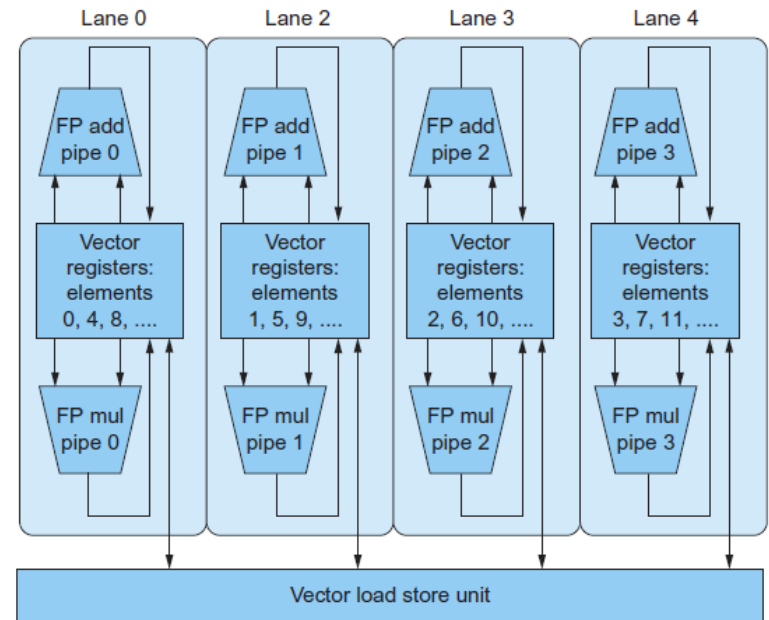
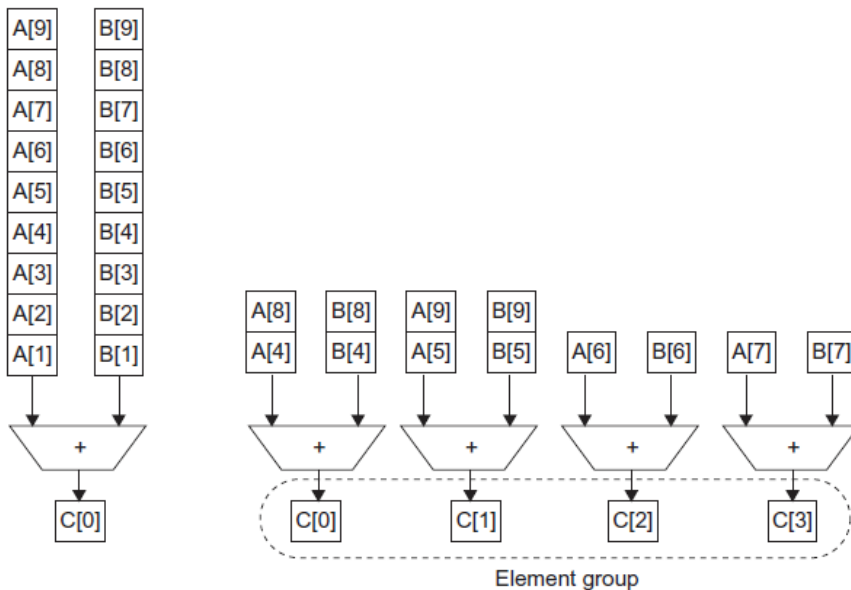
Some vector processors use registers to locate data items  
vector registers need not be loaded from adjacent locations  
These indexed loads are called a *gather* operation

Indexing via registers can also be used to store vector results  
Elements in vector registers are dispersed to non-adjacent locations  
This is known as a *scatter* operation

*Strip mining* compensates for mismatch in vector register length  
E.g. if vector is of length 60 and register is of length 32  
Vector register is loaded with first 32 elements  
And VLR (vector length register) is adjusted to load the next 28

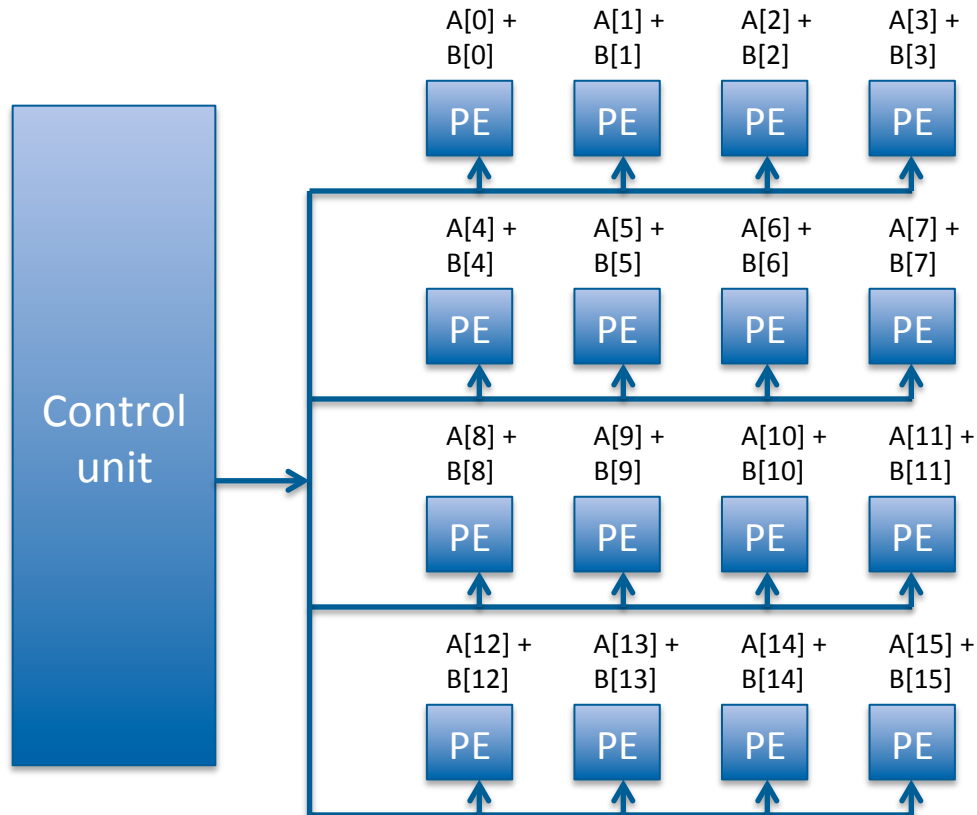


- Element  $n$  of vector register A is “hardwired” to element  $n$  of vector register B
  - Allows for multiple hardware lanes
  - No communication between lanes
  - Little increase in control overhead
  - No need to change machine code



Adding more lanes allows designers to tradeoff clock rate and energy without sacrificing performance!

- An alternative in an array processor (SIMD)
- Separate processing elements could add corresponding array elements



One control unit broadcasts the same command to multiple processing elements that operate in lock-step fashion.