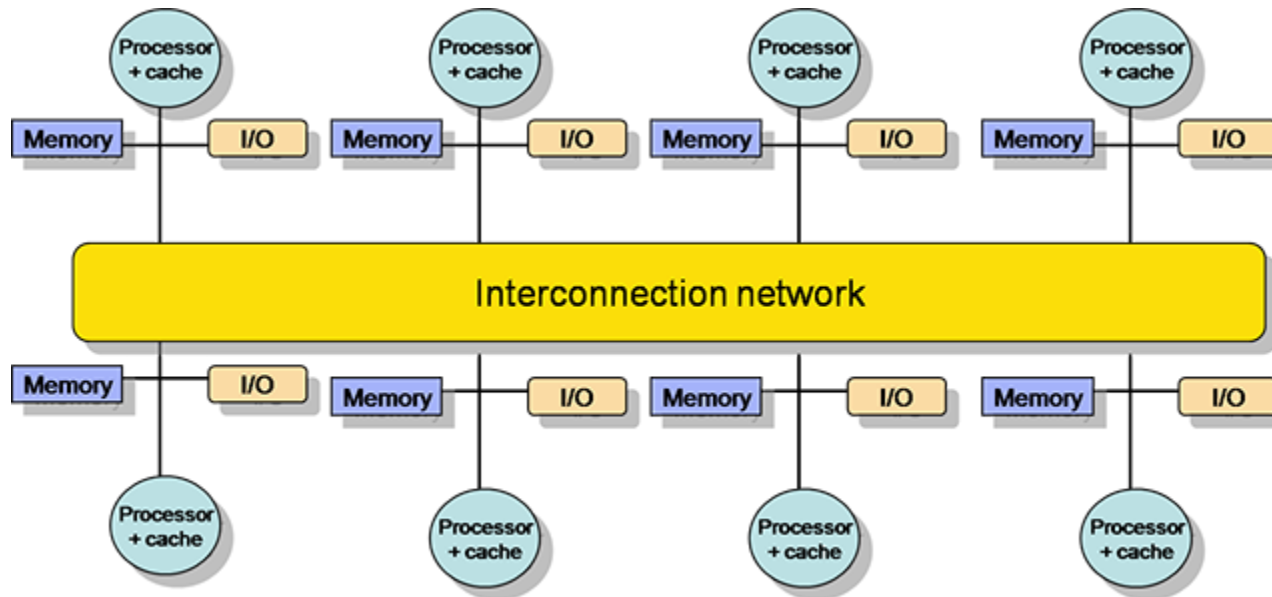



# NORMA (no remote memory access) systems



Each processor has a separate private memory & address space

# Processors only have direct access to their local memory

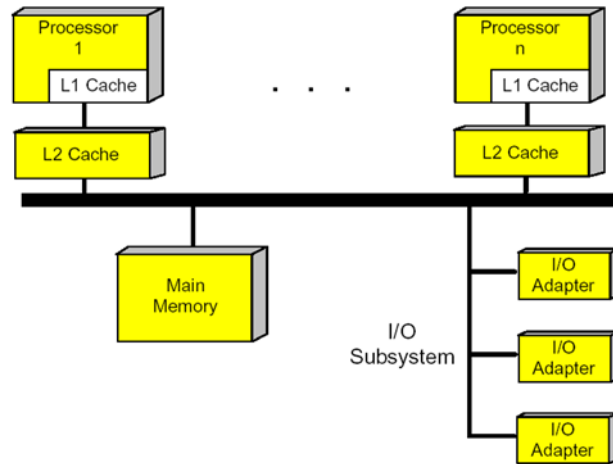


Message exchange is used to obtain data from remote modules  
Clusters fit this loosely coupled model

These are easier to scale than the SMP tightly coupled counterpart

## Using a single global memory unit

Uniform memory access (UMA) for all processors



Only one processor at a time can access memory

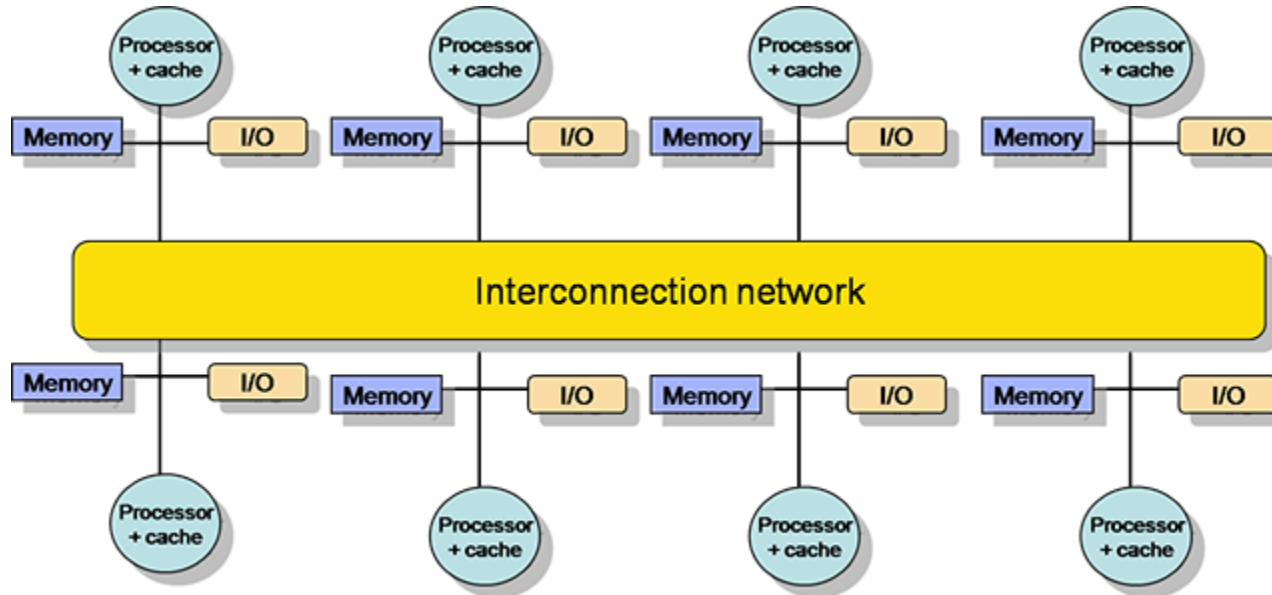
Competition for the CPU-to-memory bus will be high

Cache can be used to reduce the conflicts

This would be too limiting

# Distributed memory modules work better

## Non-uniform memory access (NUMA) system



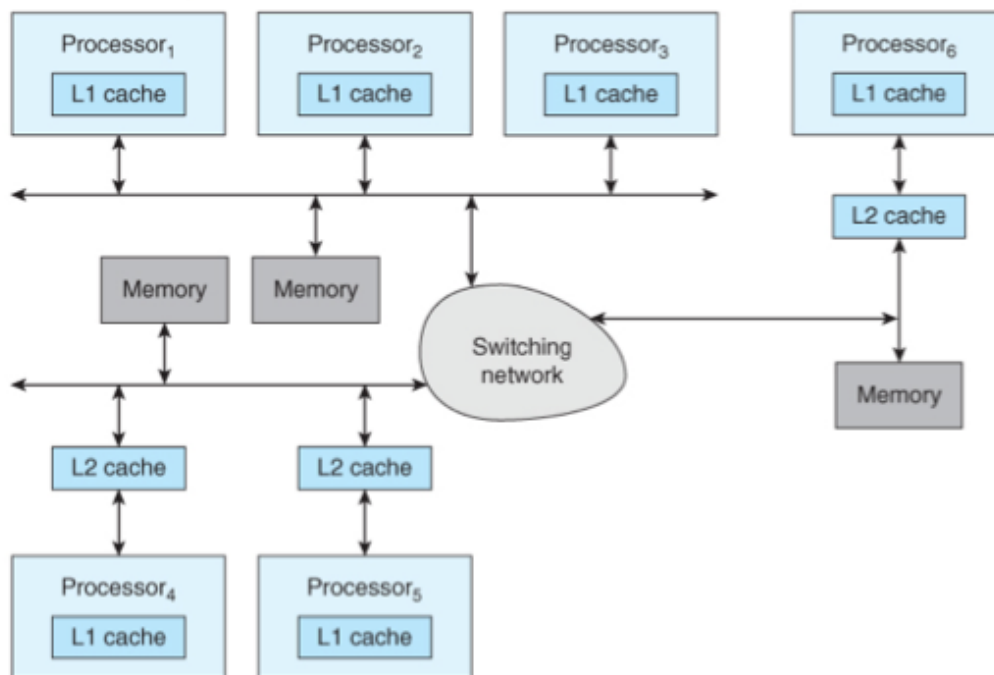
Local module access is faster than to a remote module

CPU addresses map to a particular memory module

Single shared address space spans all modules

Each processor has a cache system

Another design for a NUMA system is shown below:



Memory is accessed less often when cache is used  
But shared data must be kept consistent (i.e., coherent)

## Correct program order must be preserved

Each processor must read the most recent version of data



## All writes must be visible to all processors


If P1 writes x and then P2 reads x, P2 must obtain value written by P1

## Order must be preserved


Example: P1 sets x to 1, P2 reads x and if  $x \neq 1$ , sets it to 2

If P3 then reads x, it should see 2 and not the old value 1

a) Initially  $Q = 7$  in memoryb) P1 reads  $Q$  and caches it locally




c) P3 reads Q from memory and caches it.




d) P3 writes Q=12 to cache and to memory





e) P2 reads Q from memory and caches it



f) P2 computes  $Q + 2$ , caches it locally, and writes it back to memory

There are now three different cached values of Q  
This happened because cache coherency was not maintained.